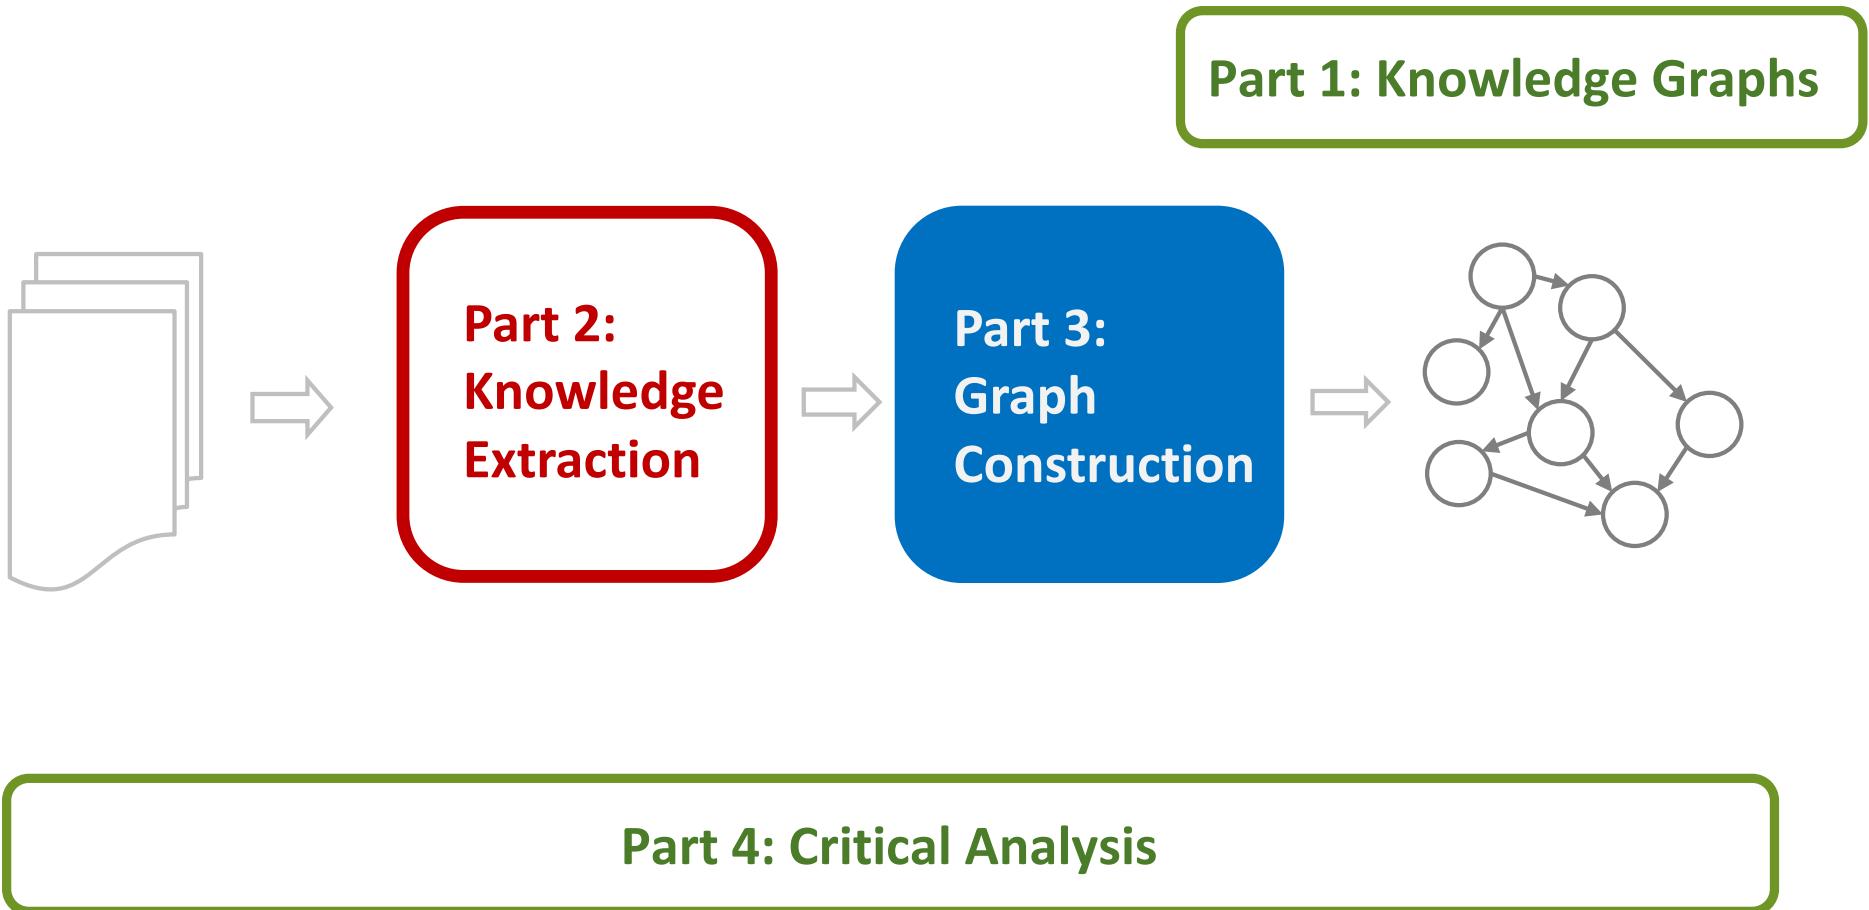


Knowledge Graph Construction from Text

AAAI 2017

JAY PUJARA, SAMEER SINGH, BHAVANA DALVI

Tutorial Overview



Tutorial Outline

1. Knowledge Graph Primer [Jay] 
2. Knowledge Extraction from Text
 - a. NLP Fundamentals [Sameer] 
 - b. Information Extraction [Bhavana] 
- Coffee Break 
3. Knowledge Graph Construction
 - a. Probabilistic Models [Jay] 
 - b. Embedding Techniques [Sameer] 
4. Critical Overview and Conclusion [Bhavana] 

Knowledge Graph Construction

TOPICS:

PROBLEM SETTING

PROBABILISTIC MODELS

EMBEDDING TECHNIQUES

Knowledge Graph Construction

TOPICS:

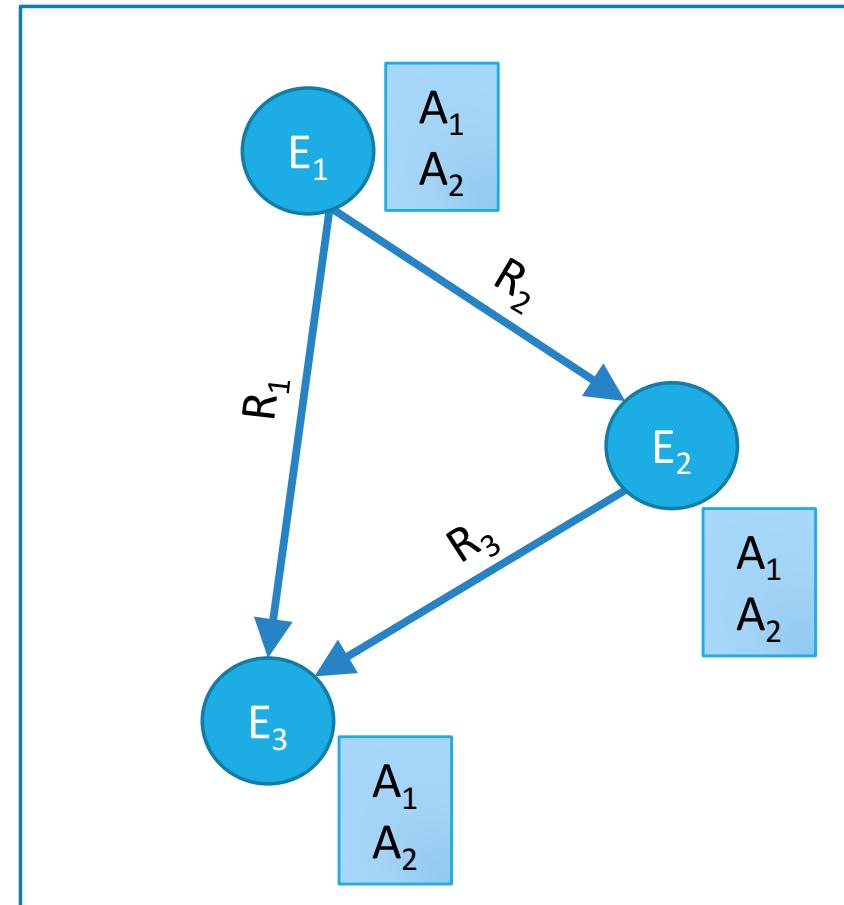
PROBLEM SETTING

PROBABILISTIC MODELS

EMBEDDING TECHNIQUES

Reminder: Basic problems

- **Who** are the entities (nodes) in the graph?
- **What** are their attributes and types (labels)?
- **How** are they related (edges)?



Graph Construction Issues

Extracted knowledge is:

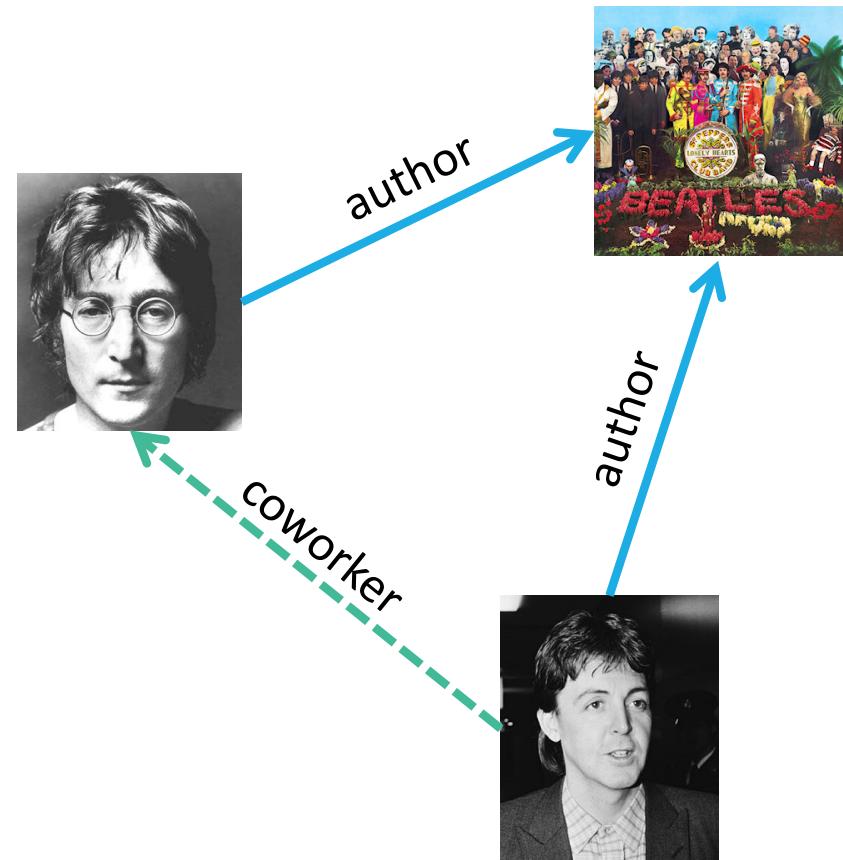
- ambiguous:
 - Ex: Beetles, beetles, Beatles
 - Ex: citizenOf, livedIn, bornIn



Graph Construction Issues

Extracted knowledge is:

- ambiguous
- incomplete
 - Ex: missing relationships
 - Ex: missing labels
 - Ex: missing entities



Graph Construction Issues

Extracted knowledge is:

- ambiguous
- incomplete
- inconsistent
 - Ex: Cynthia Lennon, Yoko Ono
 - Ex: exclusive labels (alive, dead)
 - Ex: domain-range constraints



spouse



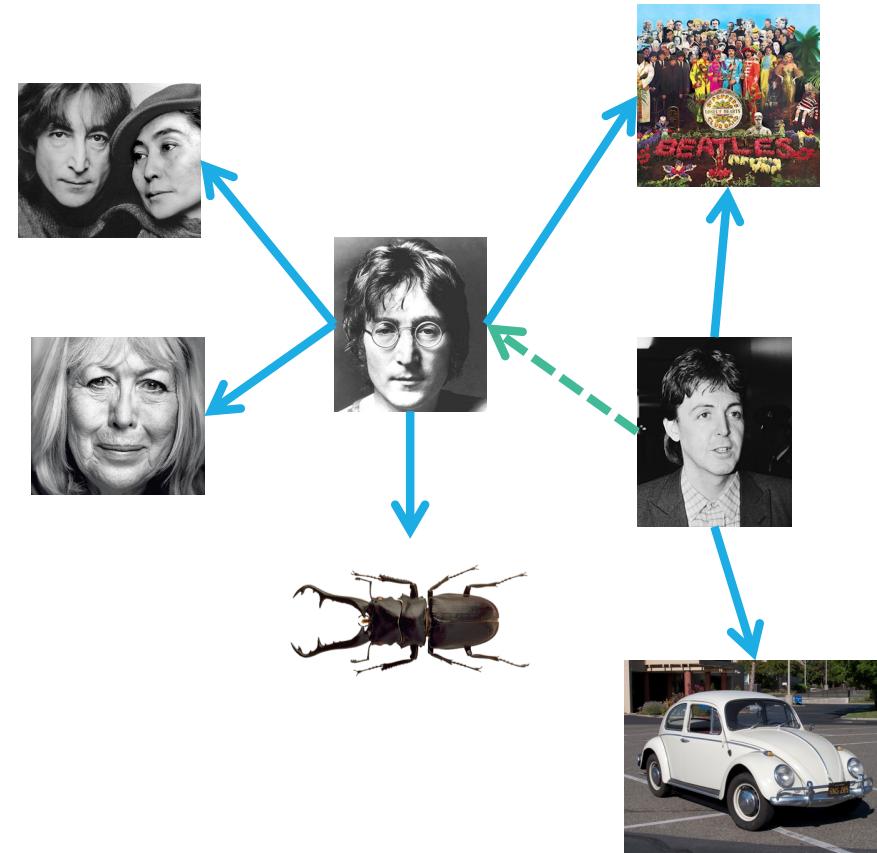
spouse



Graph Construction Issues

Extracted knowledge is:

- ambiguous



- incomplete

- inconsistent

Graph Construction approach

- Graph construction **cleans** and **completes** extraction graph
- Incorporate ontological constraints and relational patterns
- Discover statistical relationships within knowledge graph

Knowledge Graph Construction

TOPICS:

PROBLEM SETTING

PROBABILISTIC MODELS

EMBEDDING TECHNIQUES

Graph Construction Probabilistic Models

TOPICS:

OVERVIEW

GRAPHICAL MODELS

RANDOM WALK METHODS

Graph Construction Probabilistic Models

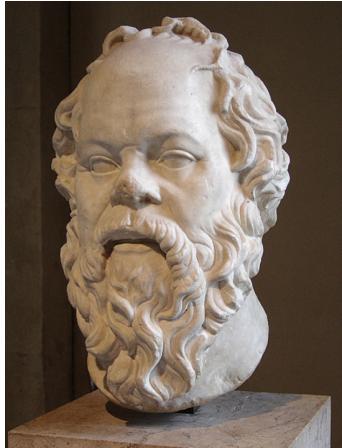
TOPICS:

OVERVIEW

GRAPHICAL MODELS

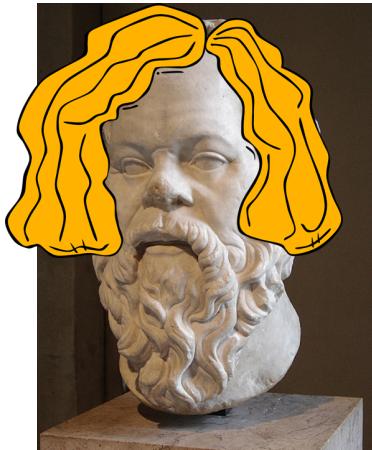
RANDOM WALK METHODS

Beyond Pure Reasoning



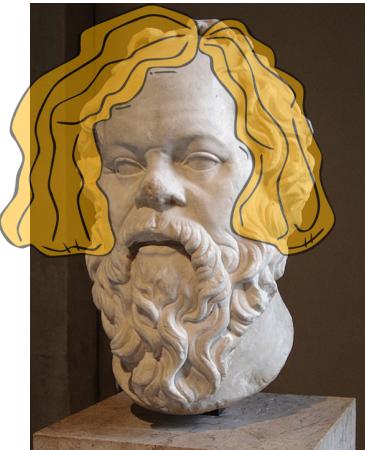
- Classical AI approach to knowledge: reasoning
 $\text{Lbl}(\text{Socrates}, \text{Man}) \ \& \ \text{Sub}(\text{Man}, \text{Mortal}) \rightarrow \text{Lbl}(\text{Socrates}, \text{Mortal})$

Beyond Pure Reasoning



- Classical AI approach to knowledge: reasoning
 $\text{Lbl}(\text{Socrates}, \text{Man}) \ \& \ \text{Sub}(\text{Man}, \text{Mortal}) \rightarrow \text{Lbl}(\text{Socrates}, \text{Mortal})$
- Reasoning difficult when extracted knowledge has errors

Beyond Pure Reasoning



- Classical AI approach to knowledge: reasoning
 $\text{Lbl}(\text{Socrates}, \text{Man}) \& \text{Sub}(\text{Man}, \text{Mortal}) \rightarrow \text{Lbl}(\text{Socrates}, \text{Mortal})$
- Reasoning difficult when extracted knowledge has errors
- Solution: probabilistic models
 $P(\text{Lbl}(\text{Socrates}, \text{Mortal}) | \text{Lbl}(\text{Socrates}, \text{Man})) = 0.9$

Graph Construction Probabilistic Models

TOPICS:

OVERVIEW

GRAPHICAL MODELS

RANDOM WALK METHODS

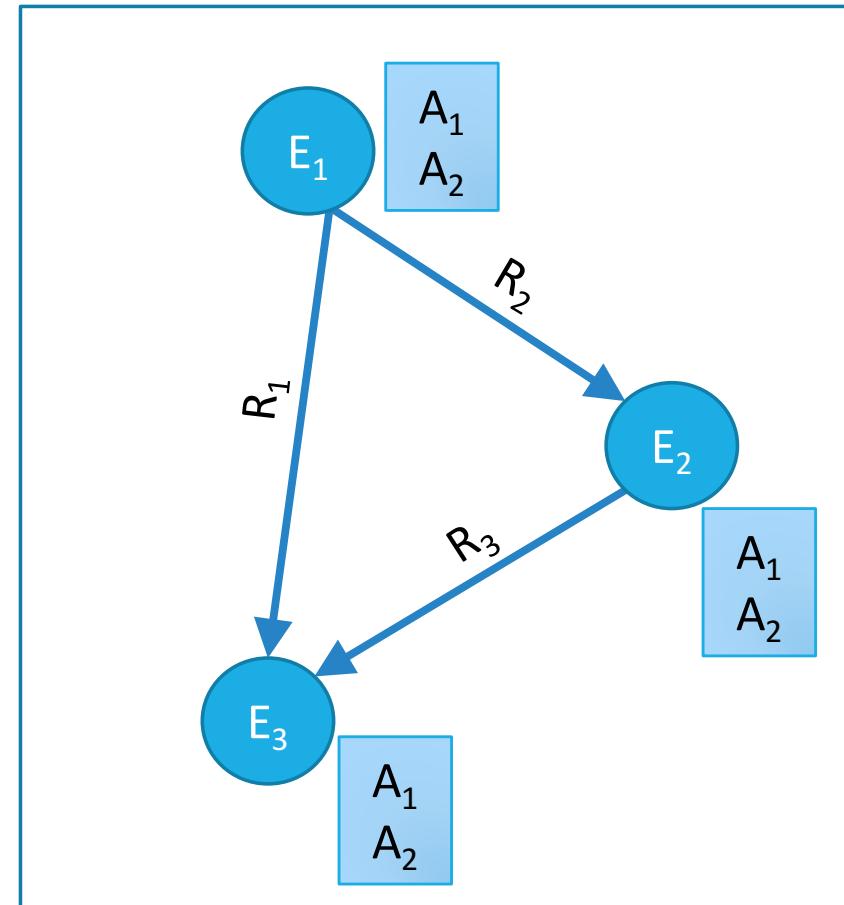
Graphical Models: Overview

- Define **joint probability distribution** on knowledge graphs
- Each candidate fact in the knowledge graph is a **variable**
- Statistical signals, ontological knowledge and rules parameterize the **dependencies** between variables
- Find most likely knowledge graph by **optimization/sampling**

Knowledge Graph Identification

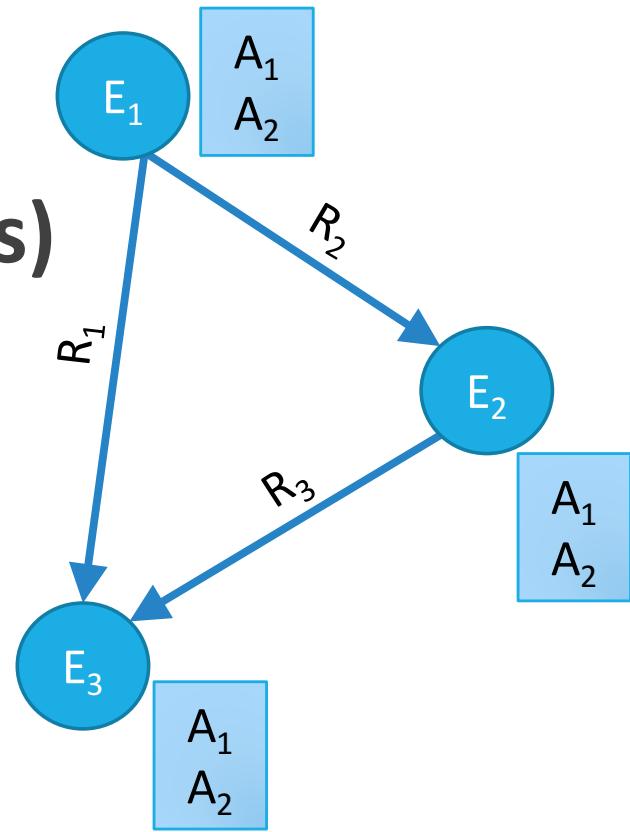
Define a graphical model to perform all three of these tasks simultaneously!

- **Who** are the entities (nodes) in the graph?
- **What** are their attributes and types (labels)?
- **How** are they related (edges)?



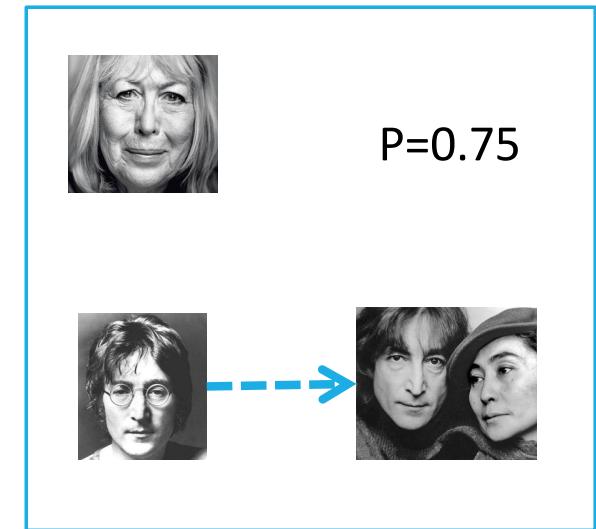
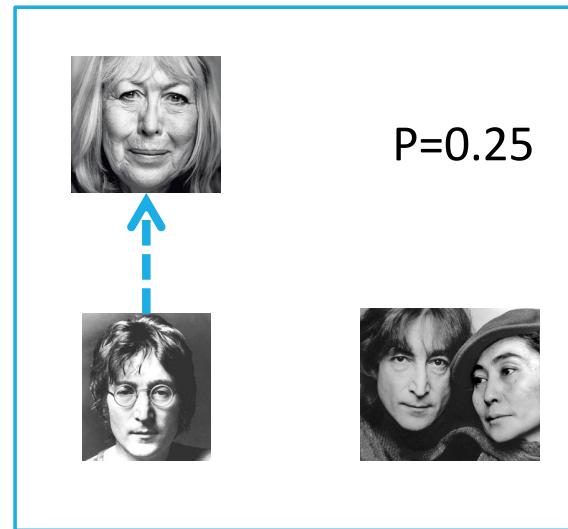
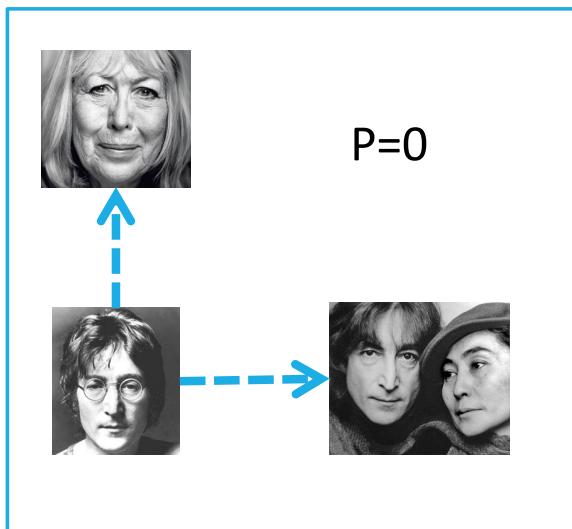
Knowledge Graph Identification

$P(\text{Who, What, How} \mid \text{Extractions})$



Probabilistic Models

- Use dependencies between facts in KG
- Probability defined *jointly* over facts



What determines probability?

- Statistical signals from text extractors and classifiers

What determines probability?

- **Statistical signals from text extractors and classifiers**
 - $P(R(\text{John}, \text{Spouse}, \text{Yoko}))=0.75$; $P(R(\text{John}, \text{Spouse}, \text{Cynthia}))=0.25$
 - LevenshteinSimilarity(Beatles, Beetles) = 0.9

What determines probability?

- Statistical signals from text extractors and classifiers
- Ontological knowledge about domain

What determines probability?

- Statistical signals from text extractors and classifiers
- Ontological knowledge about domain
 - Functional(Spouse) & R(A,Spouse,B) -> !R(A,Spouse,C)
 - Range(Spouse, Person) & R(A,Spouse,B) -> Type(B, Person)

What determines probability?

- Statistical signals from text extractors and classifiers
- Ontological knowledge about domain
- Rules and patterns mined from data

What determines probability?

- Statistical signals from text extractors and classifiers
- Ontological knowledge about domain
- Rules and patterns mined from data
 - $R(A, \text{Spouse}, B) \& R(A, \text{Lives}, L) \rightarrow R(B, \text{Lives}, L)$
 - $R(A, \text{Spouse}, B) \& R(A, \text{Child}, C) \rightarrow R(B, \text{Child}, C)$

What determines probability?

- **Statistical signals from text extractors and classifiers**
 - $P(R(\text{John}, \text{Spouse}, \text{Yoko}))=0.75$; $P(R(\text{John}, \text{Spouse}, \text{Cynthia}))=0.25$
 - LevenshteinSimilarity(Beatles, Beetles) = 0.9
- **Ontological knowledge about domain**
 - Functional(Spouse) & $R(A, \text{Spouse}, B) \rightarrow !R(A, \text{Spouse}, C)$
 - Range(Spouse, Person) & $R(A, \text{Spouse}, B) \rightarrow \text{Type}(B, \text{Person})$
- **Rules and patterns mined from data**
 - $R(A, \text{Spouse}, B) \& R(A, \text{Lives}, L) \rightarrow R(B, \text{Lives}, L)$
 - $R(A, \text{Spouse}, B) \& R(A, \text{Child}, C) \rightarrow R(B, \text{Child}, C)$

Example: The Fab Four

THE
BEATLES



Illustration of KG Identification

Uncertain Extractions:

- .5: Lbl(Fab Four, novel)
- .7: Lbl(Fab Four, musician)
- .9: Lbl(Beatles, musician)
- .8: Rel(Beatles, AlbumArtist,
Abbey Road)

Illustration of KG Identification

Uncertain Extractions:

- .5: Lbl(Fab Four, novel)
- .7: Lbl(Fab Four, musician)
- .9: Lbl(Beatles, musician)
- .8: Rel(Beatles, AlbumArtist, Abbey Road)

(Annotated) Extraction Graph

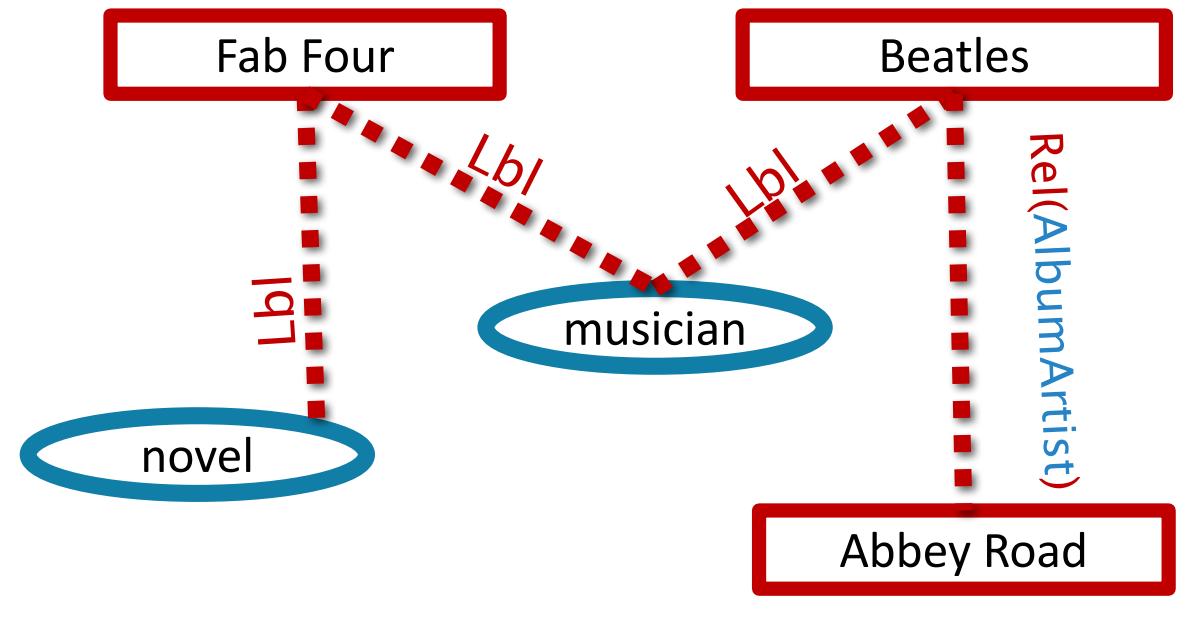


Illustration of KG Identification

Uncertain Extractions:

- .5: Lbl(Fab Four, novel)
- .7: Lbl(Fab Four, musician)
- .9: Lbl(Beatles, musician)
- .8: Rel(Beatles, AlbumArtist,
 Abbey Road)

Ontology:

- Dom(albumArtist, musician)
- Mut(novel, musician)

Extraction Graph

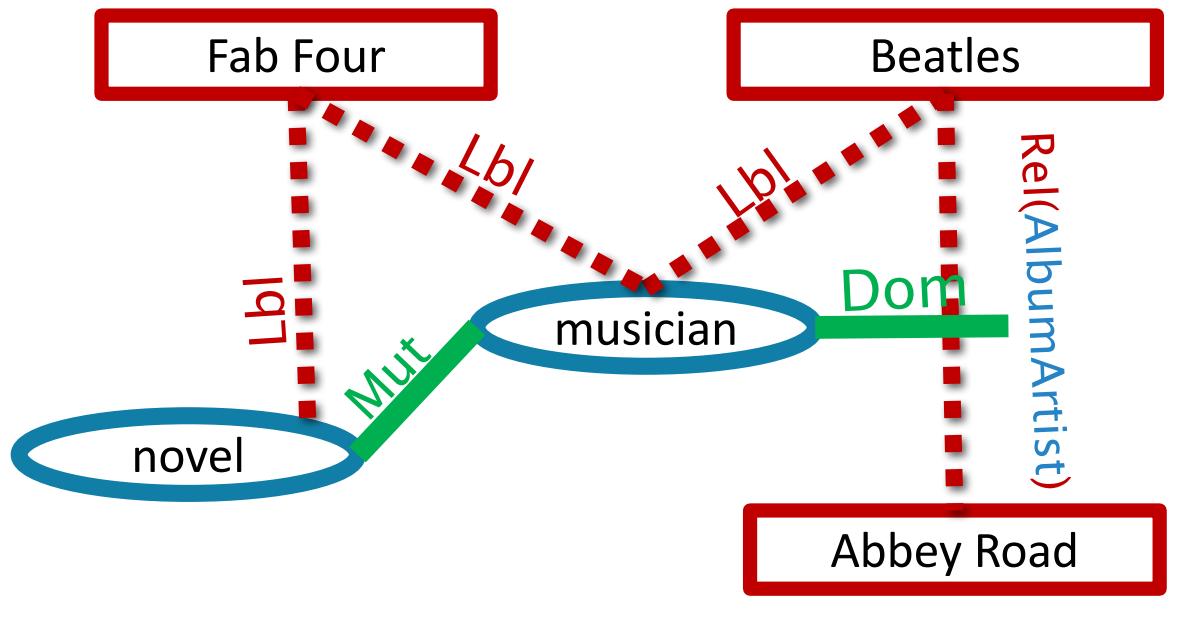


Illustration of KG Identification

Uncertain Extractions:

- .5: Lbl(Fab Four, novel)
- .7: Lbl(Fab Four, musician)
- .9: Lbl(Beatles, musician)
- .8: Rel(Beatles, AlbumArtist, Abbey Road)

Ontology:

- Dom(albumArtist, musician)
- Mut(novel, musician)

Entity Resolution:

SameEnt(Fab Four, Beatles)

(Annotated) Extraction Graph

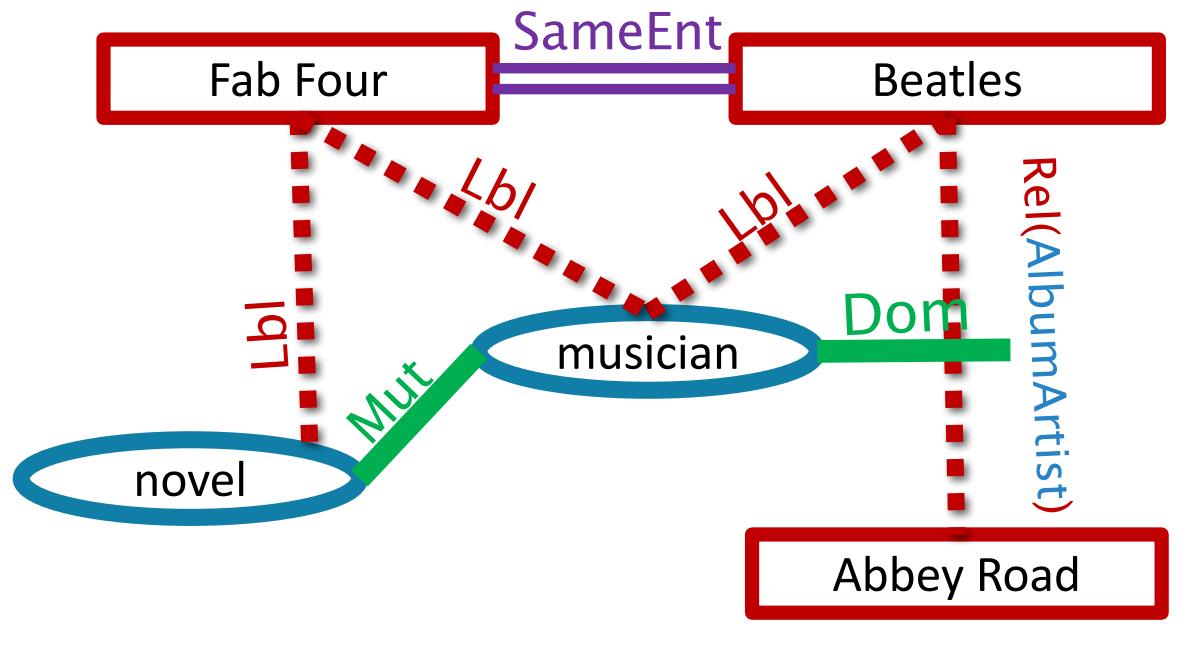


Illustration of KG Identification

Uncertain Extractions:

- .5: Lbl(Fab Four, novel)
- .7: Lbl(Fab Four, musician)
- .9: Lbl(Beatles, musician)
- .8: Rel(Beatles, AlbumArtist, Abbey Road)

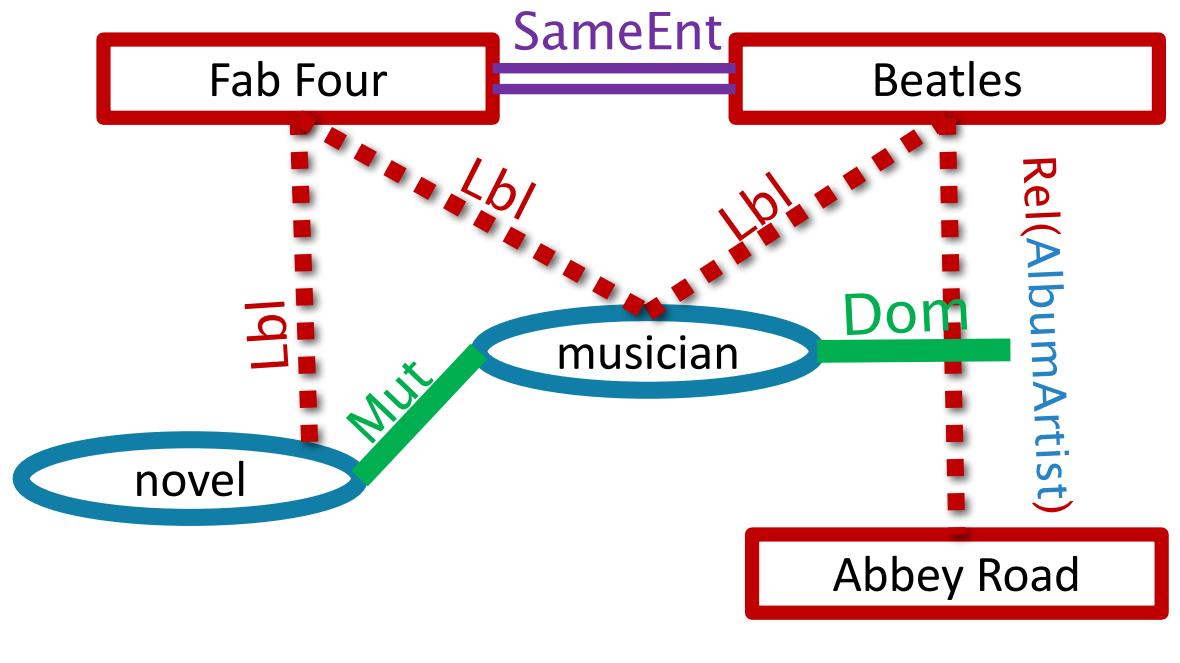
Ontology:

- Dom(albumArtist, musician)
- Mut(novel, musician)

Entity Resolution:

SameEnt(Fab Four, Beatles)

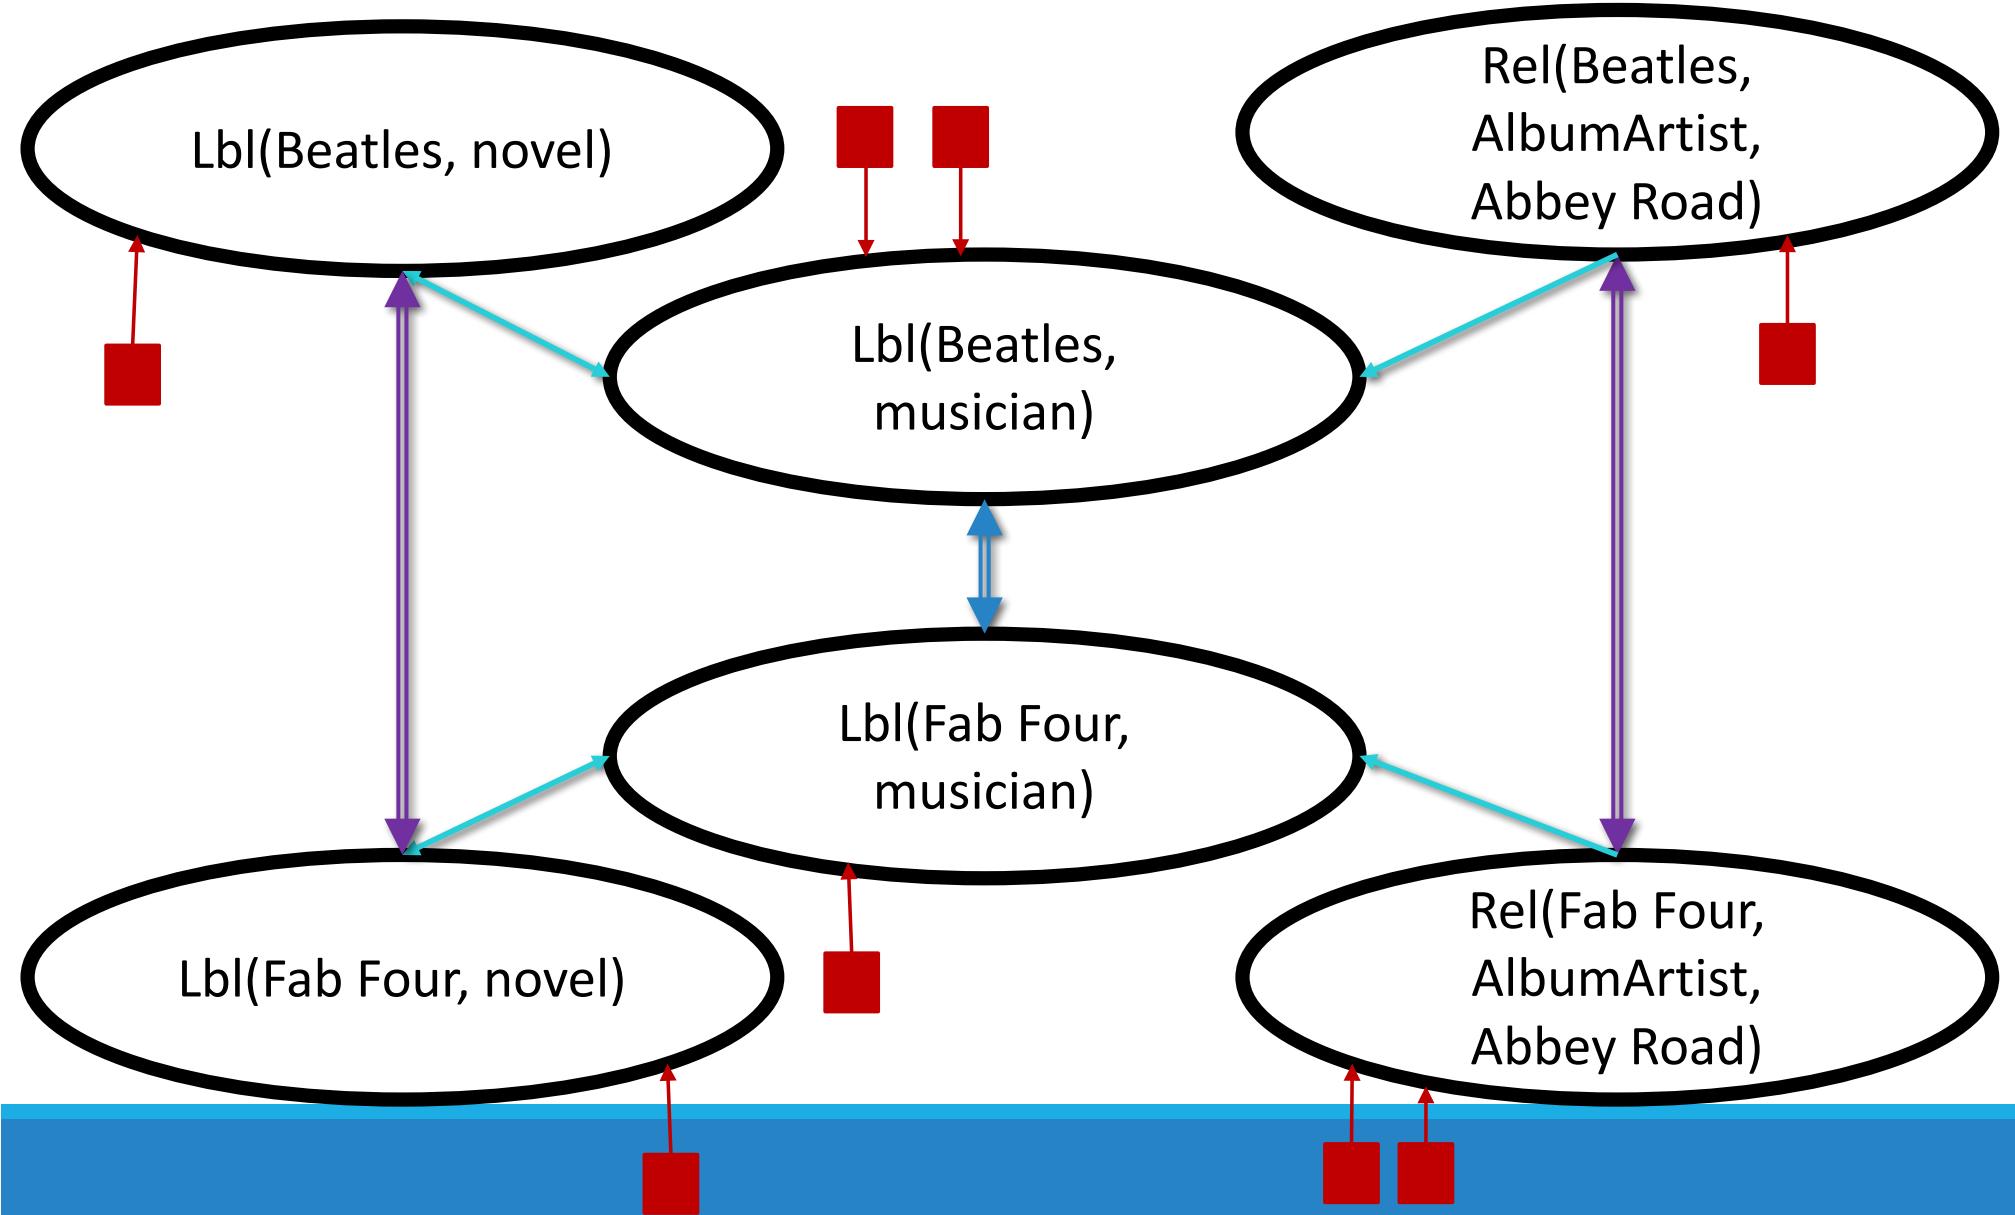
(Annotated) Extraction Graph



After Knowledge Graph Identification



Probabilistic graphical model for KG



Defining graphical models

- Many options for defining a graphical model
- We focus on two approaches, MLNs and PSL, that use **rules**
- **MLNs** treat facts as Boolean, use sampling for satisfaction
- **PSL** infers a “truth value” for each fact via optimization



Rules for KG Model

100:	Subsumes(L1,L2)	& Label(E,L1)	->	Label(E,L2)
100:	Exclusive(L1,L2)	& Label(E,L1)	->	!Label(E,L2)
100:	Inverse(R1,R2)	& Relation(R1,E,0)	->	Relation(R2,0,E)
100:	Subsumes(R1,R2)	& Relation(R1,E,0)	->	Relation(R2,E,0)
100:	Exclusive(R1,R2)	& Relation(R1,E,0)	->	!Relation(R2,E,0)
100:	Domain(R,L)	& Relation(R,E,0)	->	Label(E,L)
100:	Range(R,L)	& Relation(R,E,0)	->	Label(0,L)
10:	SameEntity(E1,E2)	& Label(E1,L)	->	Label(E2,L)
10:	SameEntity(E1,E2)	& Relation(R,E1,0)	->	Relation(R,E2,0)
1:	Label_NYT(E,L)		->	Label(E,L)
1:	Label_YouTube(E,L)		->	Label(E,L)
1:	Relation_LATimes(R,E,0)		->	Relation(R,E,0)
1:			->	!Relation(R,E,0)
1:			->	!Label(E,L)

Rules to Distributions

- Rules are *grounded* by substituting literals into formulas

$w_r : \text{SAMEENT}(\text{Fab Four}, \text{Beatles}) \wedge$

$\text{LBL}(\text{Beatles}, \text{musician}) \Rightarrow \text{LBL}(\text{Fab Four}, \text{musician})$

- Each ground rule has a weighted satisfaction derived from the formula's truth value

$$P(G|E) = \frac{1}{Z} \exp \left[\sum_{r \in R} w_r \phi_r(G, E) \right]$$

- Together, the ground rules provide a joint probability distribution over knowledge graph facts, conditioned on the extractions

Probability Distribution over KGs

$$P(G | E) = \frac{1}{Z} \exp \left[- \sum_{r \in R} w_r \varphi_r(G) \right]$$

CANDLBL_T(FabFour, novel)

\Rightarrow LBL(FabFour, novel)

MUT(novel, musician)

\wedge LBL(Beatles, novel)

\Rightarrow \neg LBL(Beatles, musician)

SAMEENT(Beatles, FabFour)

\wedge LBL(Beatles, musician)

\Rightarrow LBL(FabFour, musician)

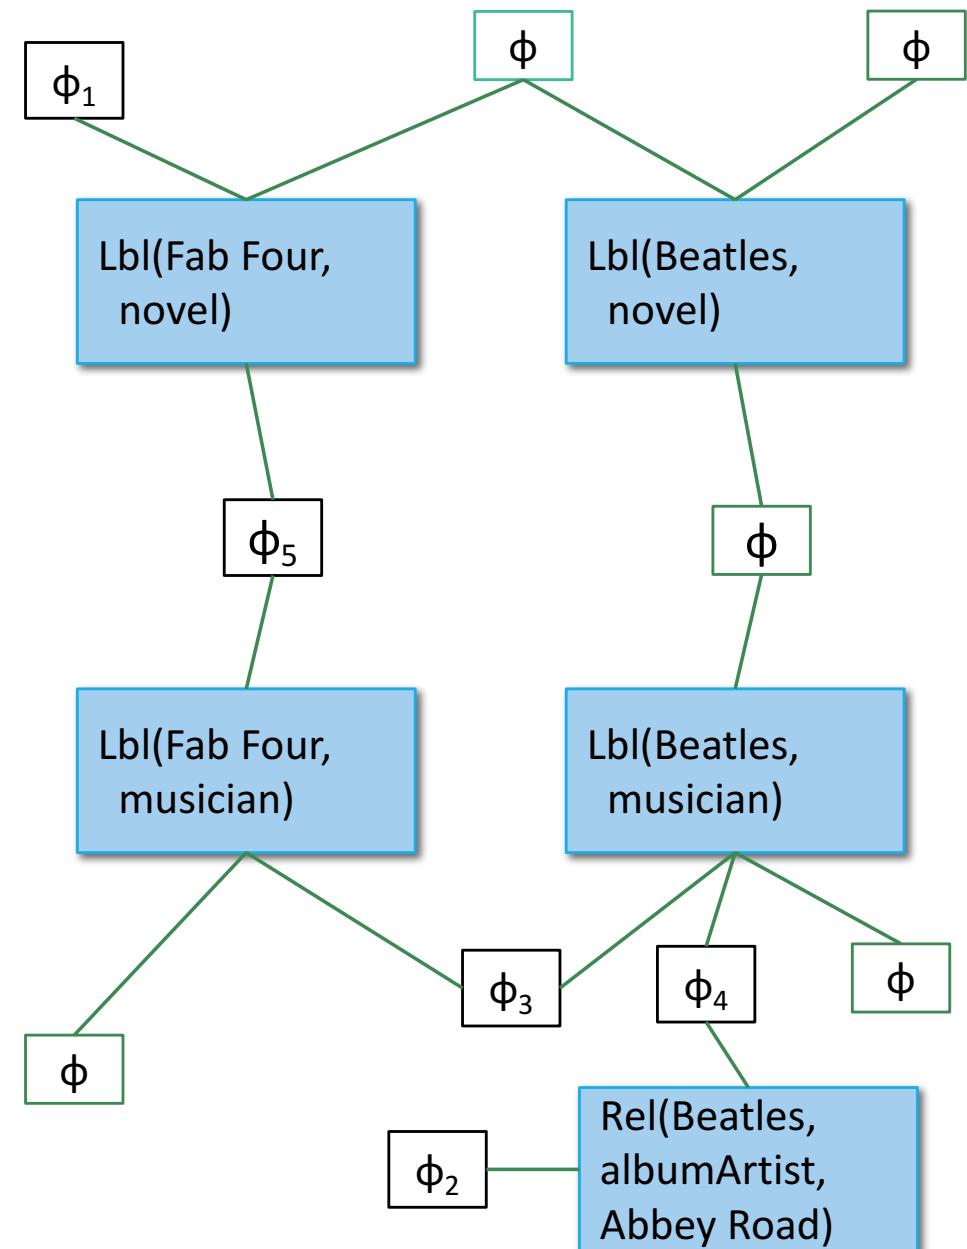
$[\phi_1] \text{ CANDLBL}_{\text{struct}}(\text{FabFour}, \text{novel})$
 $\Rightarrow \text{LBL}(\text{FabFour}, \text{novel})$

$[\phi_2] \text{ CANDREL}_{\text{pat}}(\text{Beatles}, \text{AlbumArtist}, \text{AbbeyRoad})$
 $\Rightarrow \text{REL}(\text{Beatles}, \text{AlbumArtist}, \text{AbbeyRoad})$

$[\phi_3] \text{ SAMEENT}(\text{Beatles}, \text{FabFour})$
 $\wedge \text{LBL}(\text{Beatles}, \text{musician})$
 $\Rightarrow \text{LBL}(\text{FabFour}, \text{musician})$

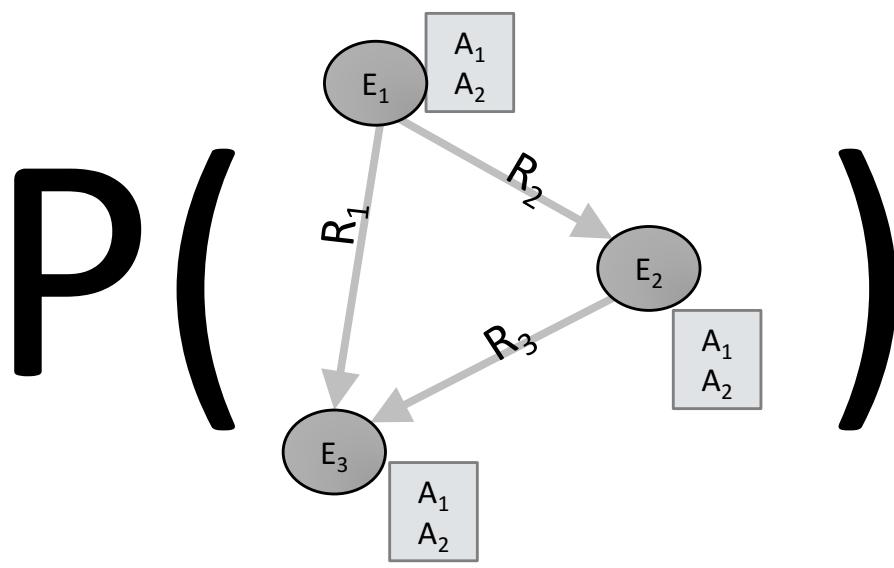
$[\phi_4] \text{ DOM}(\text{AlbumArtist}, \text{musician})$
 $\wedge \text{REL}(\text{Beatles}, \text{AlbumArtist}, \text{AbbeyRoad})$
 $\Rightarrow \text{LBL}(\text{Beatles}, \text{musician})$

$[\phi_5] \text{ MUT}(\text{musician}, \text{novel})$
 $\wedge \text{LBL}(\text{FabFour}, \text{musician})$
 $\Rightarrow \neg \text{LBL}(\text{FabFour}, \text{novel})$

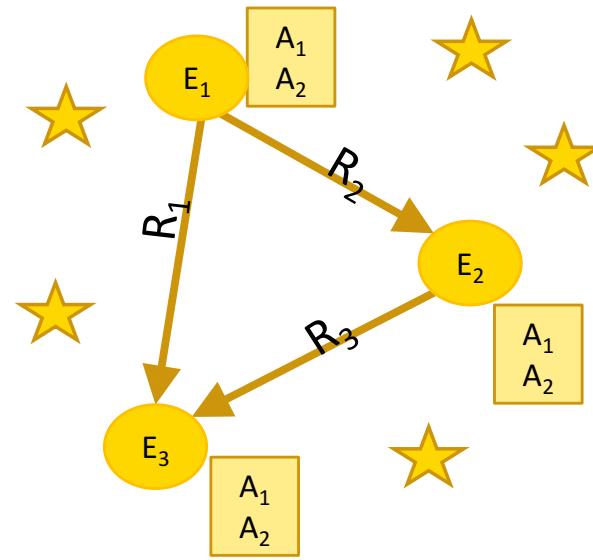


How do we get a knowledge graph?

Have: $P(KG)$ for all KGs



Need: best KG



MAP inference: optimizing over distribution to find the best knowledge graph

Inference and KG optimization

- Finding the best KG satisfying weighed rules: NP Hard
- MLNs [discrete]: Monte Carlo sampling methods
 - Solution quality dependent on burn-in time, iterations, etc.
- PSL [continuous]: optimize convex linear surrogate
 - Fast optimization, $\frac{3}{4}$ -optimal MAX SAT lower bound

Graphical Models Experiments

Data: ~1.5M extractions, ~70K ontological relations, ~500 relation/label types

Task: Collectively construct a KG and evaluate on 25K target facts

Comparisons:

Extract	Average confidences of extractors for each fact in the NELL candidates
Rules	Default, rule-based heuristic strategy used by the NELL project
MLN	Jiang+, ICDM12 – estimates marginal probabilities with MC-SAT
PSL	Pujara+, ISWC13 – convex optimization of continuous truth values with ADMM

Running Time: Inference completes in 10 seconds, values for 25K facts

	AUC	F1
Extract	.873	.828
Rules	.765	.673
MLN (Jiang, 12)	.899	.836
PSL (Pujara, 13)	.904	.853

Graphical Models: Pros/Cons

BENEFITS

- Define probability distribution over KGs
- Easily specified via rules
- Fuse knowledge from many different sources

DRAWBACKS

- Requires optimization over all KG facts - overkill
- Dependent on rules from ontology/expert
- Require probabilistic semantics - unavailable

Graph Construction Probabilistic Models

TOPICS:

OVERVIEW

GRAPHICAL MODELS

RANDOM WALK METHODS

Random Walk Overview

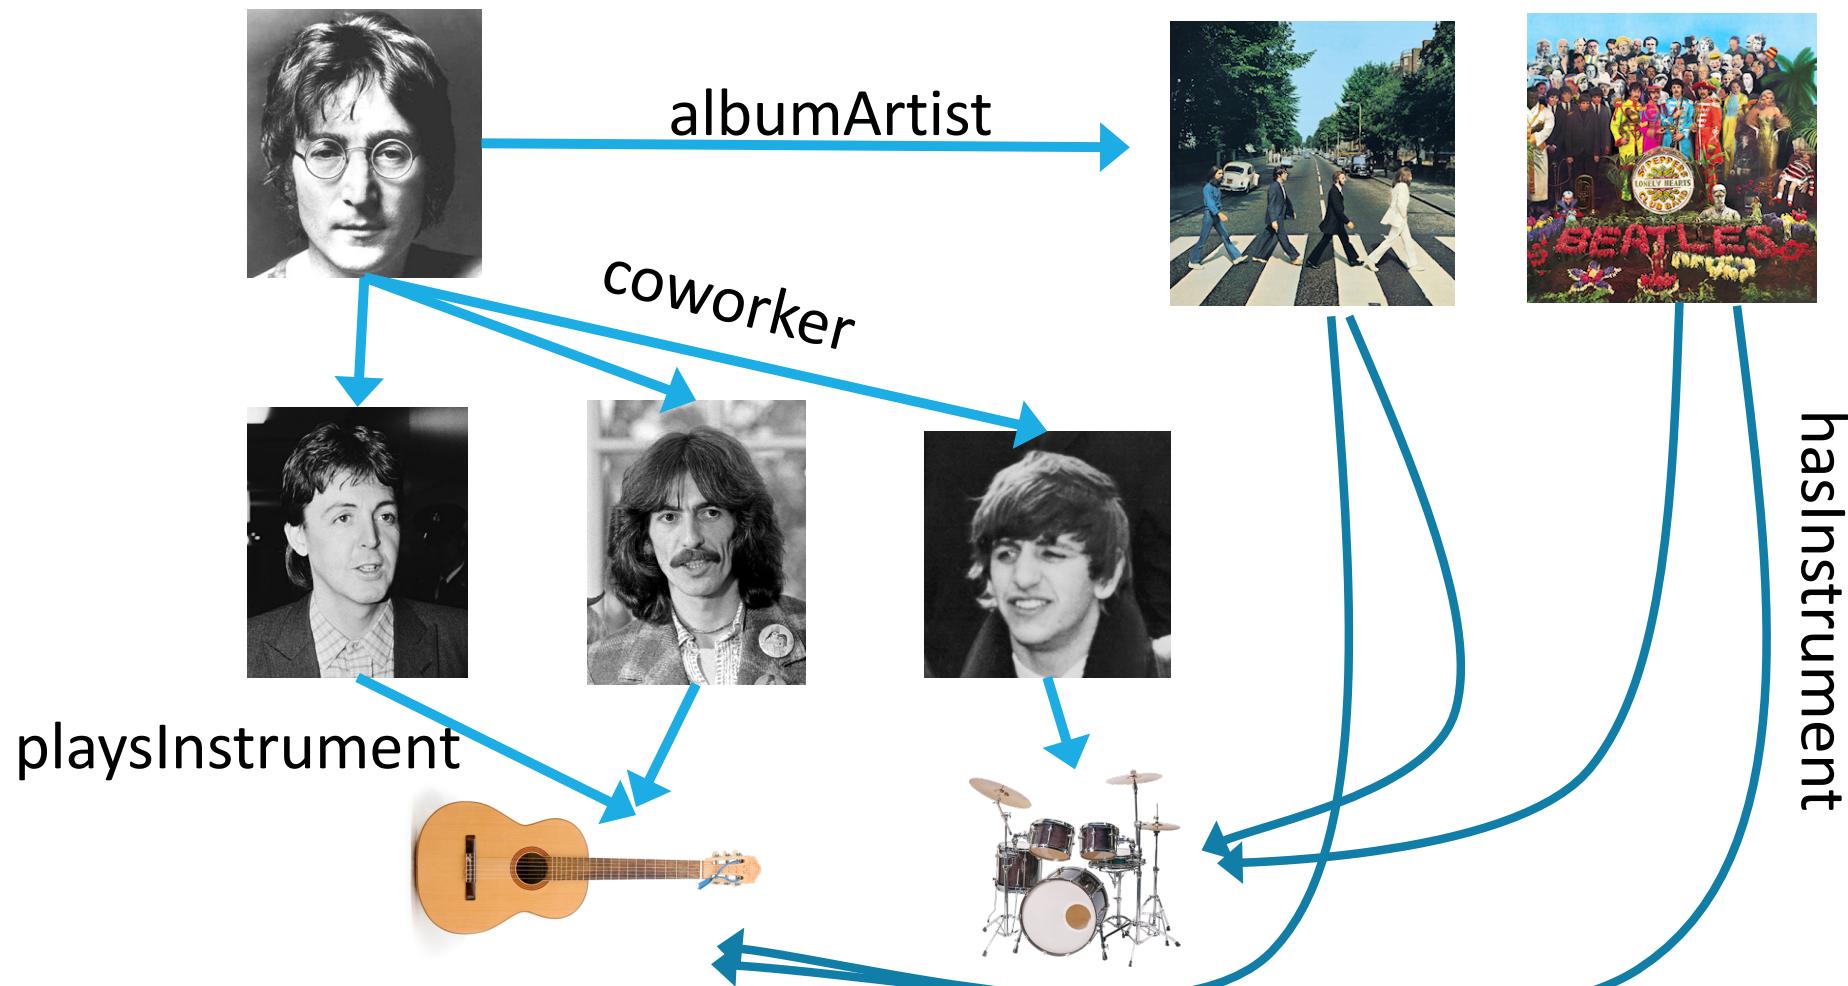
- Given: a query of an **entity** and **relation**
- Starting at the entity, **randomly walk** the KG
- Random walk ends when reaching an appropriate **goal**
- Learned **parameters** bias choices in the random walk
- Output **relative probabilities** of goal states

Random Walk Illustration

Query: R(Lennon, PlaysInstrument, ?)

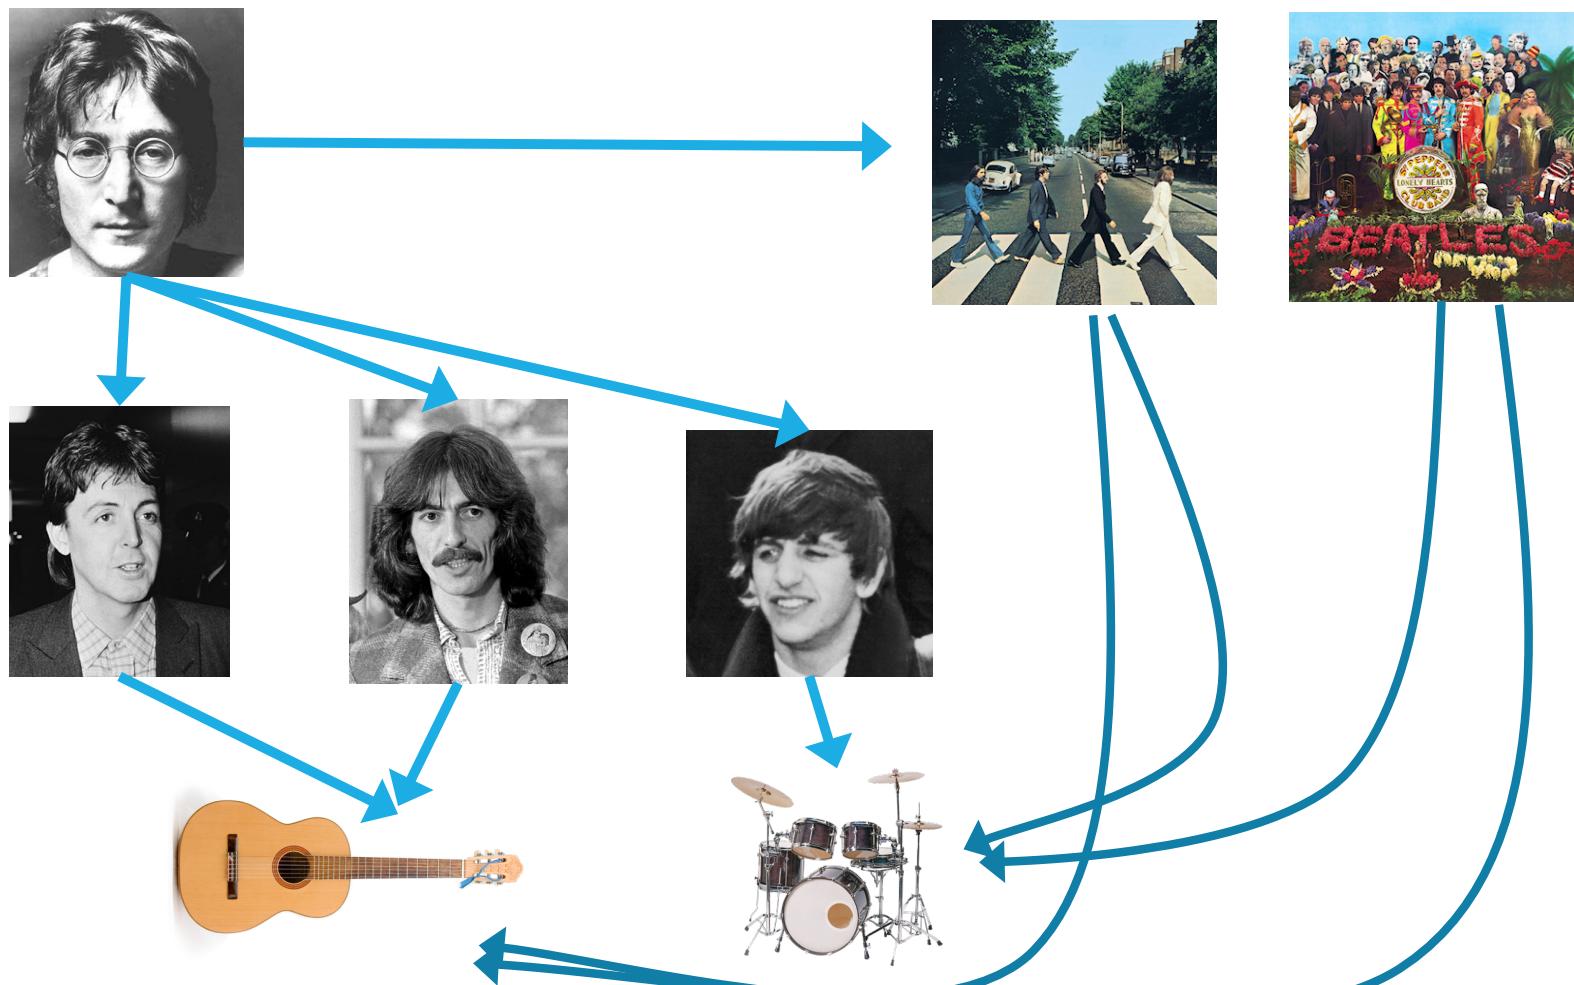
Random Walk Illustration

Query: R(Lennon, PlaysInstrument, ?)



Random Walk Illustration

Query: R(Lennon, PlaysInstrument, ?)



Random Walk Illustration

Query Q: R(Lennon, PlaysInstrument, ?)



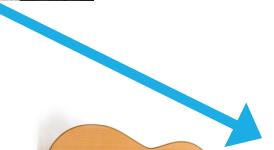
Random Walk Illustration

Query Q: R(Lennon, PlaysInstrument, ?)



Random Walk Illustration

Query Q: R(Lennon, PlaysInstrument, ?)



Random Walk Illustration

Query Q: R(Lennon, PlaysInstrument, ?)



$P(Q | \pi = \langle \text{coworker}, \text{playsInstrument} \rangle) W_\pi$

Path

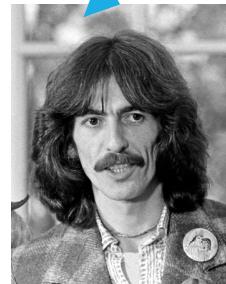
Weight of path

Random Walk Illustration

Query Q: R(Lennon, PlaysInstrument, ?)



$P(Q | \pi = \langle \text{coworker}, \text{playsInstrument} \rangle) W_\pi$



Random Walk Illustration

Query Q: R(Lennon, PlaysInstrument, ?)

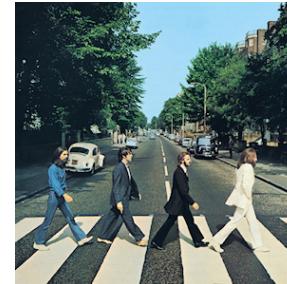


$P(Q | \pi = \langle \text{coworker}, \text{playsInstrument} \rangle) W_\pi$



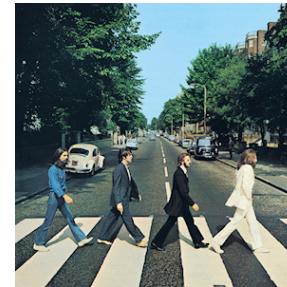
Random Walk Illustration

Query Q: R(Lennon, PlaysInstrument, ?)

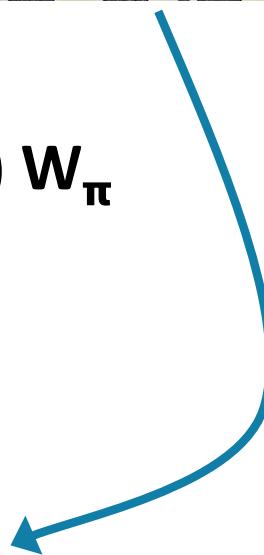


Random Walk Illustration

Query Q: R(Lennon, PlaysInstrument, ?)

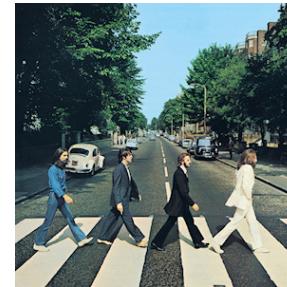


$P(Q | \pi = \langle \text{albumArtist}, \text{hasInstrument} \rangle) W_\pi$



Random Walk Illustration

Query Q: R(Lennon, PlaysInstrument, ?)

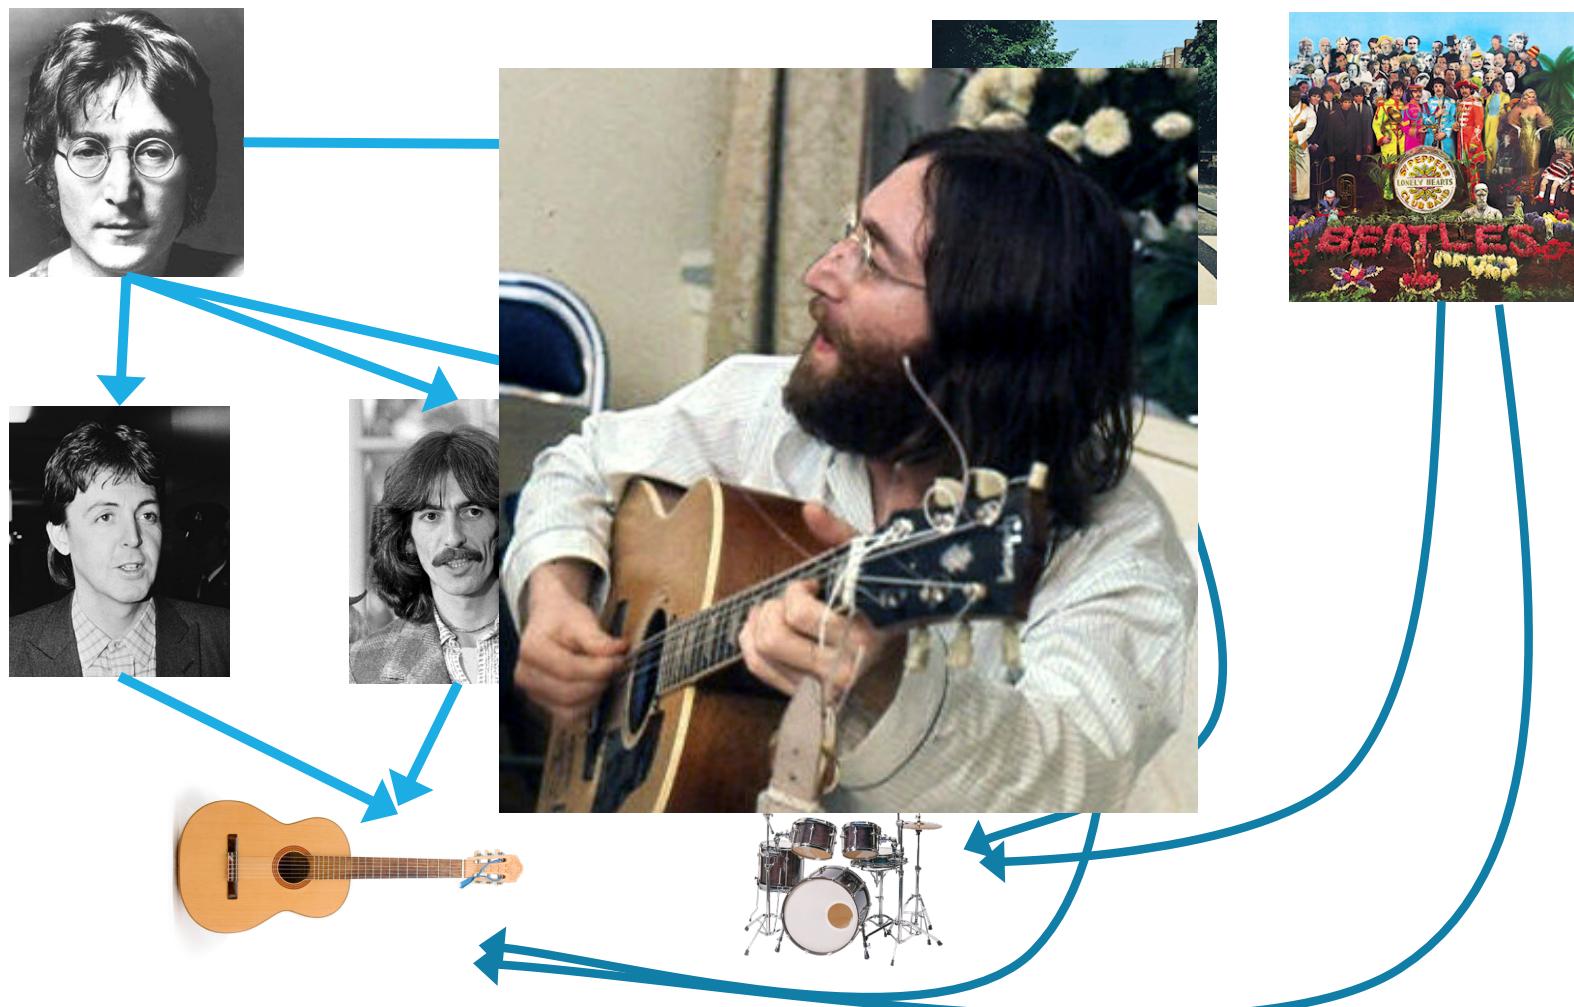


$P(Q | \pi = \text{<albumArtist, hasInstrument>}) W_\pi$



Random Walk Illustration

Query: R(Lennon, PlaysInstrument, ?)



Recent Random Walk Methods

PRA: Path Ranking Algorithm

- Performs random walk of **imperfect knowledge graph**
- Estimates **transition probabilities** using KG
- For each relation, learns **parameters for paths** through the KG

ProPPR: Programming with Personalized PageRank

- Constructs **proof graph**
 - Nodes are partially-ground clauses with one or more facts
 - Edges are proof-transformations
- **Parameters** are learned for each **ground entity and rule**

Recent Random Walk Methods

PRA: Path Ranking Algorithm

- Performs random walk of **imperfect knowledge graph**
- Estimates **transition probabilities** using KG
- For each relation, learns **parameters for paths** through the KG

ProPPR: Programming with Personalized PageRank

- Constructs **proof graph**
 - Nodes are partially-ground clauses with one or more facts
 - Edges are proof-transformations
- **Parameters** are learned for each **ground entity and rule**

PRA in a nutshell

$$\text{score}(q.s \rightarrow e; q) = \sum_{\pi_i \in \Pi_b} P(q.s \rightarrow e; \pi_i) W_{\pi_i}$$

PRA in a nutshell

$$\text{score}(q.s \rightarrow e; q) = \sum_{\pi_i \in \Pi_b} P(q.s \rightarrow e; \pi_i) W_{\pi_i}$$

Filter paths based on HITS and accuracy

PRA in a nutshell

$$\text{score}(q.s \rightarrow e; q) = \sum_{\pi_i \in \Pi_b} P(q.s \rightarrow e; \pi_i) W_{\pi_i}$$

Filter paths based on HITS and accuracy

Estimate probabilities efficiently with dynamic programming

PRA in a nutshell

$$\text{score}(q.s \rightarrow e; q) = \sum_{\pi_i \in \Pi_b} P(q.s \rightarrow e; \pi_i) W_{\pi_i}$$

Filter paths based on HITS and accuracy

Estimate probabilities efficiently with dynamic programming

Path weights are learned with logistic regression

Recent Random Walk Methods

PRA: Path Ranking Algorithm

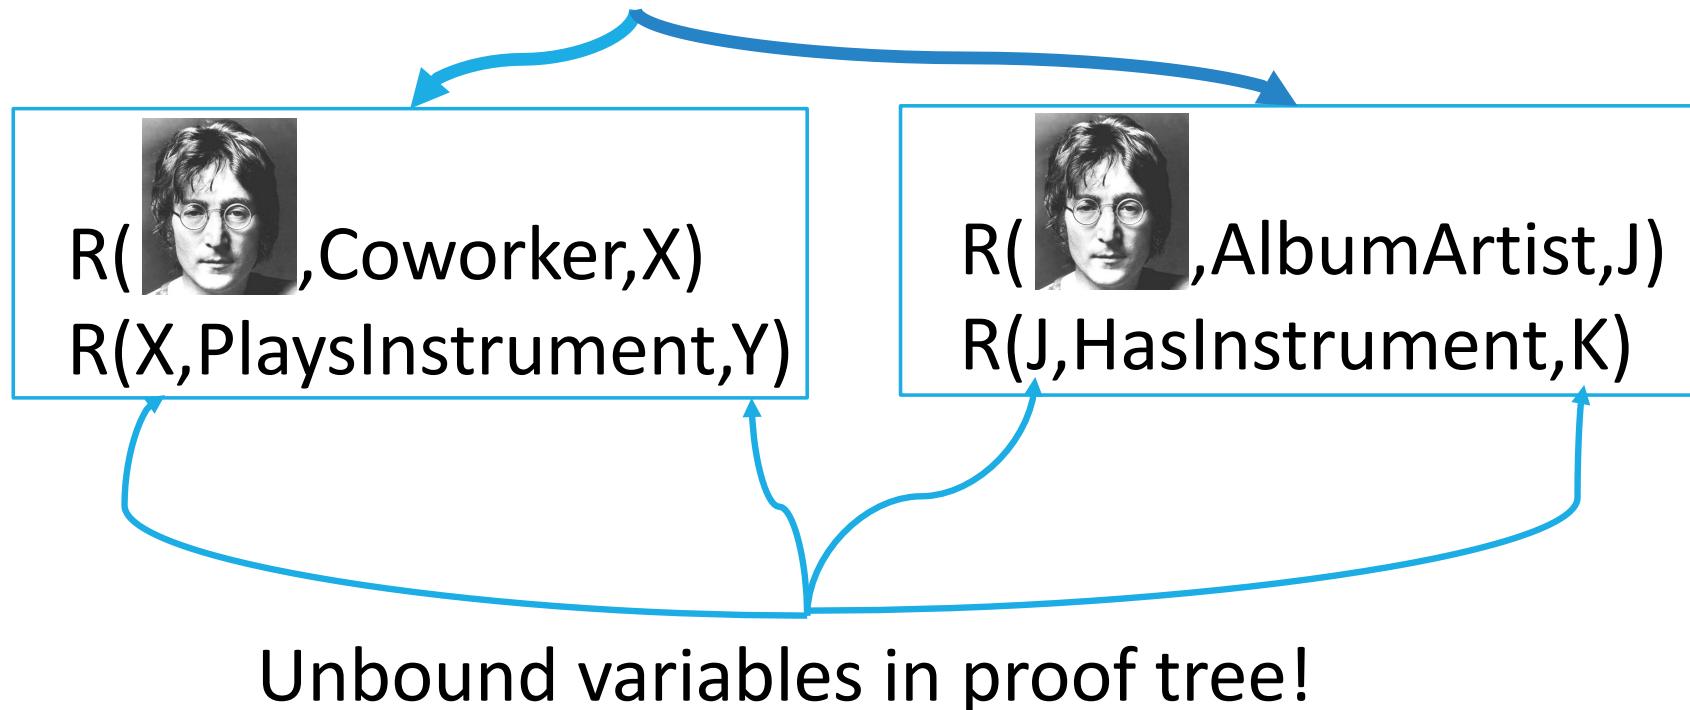
- Performs random walk of **imperfect knowledge graph**
- Estimates **transition probabilities** using KG
- For each relation, learns **parameters for paths** through the KG

ProPPR: ProbLog + Personalized PageRank

- Constructs **proof graph**
 - Nodes are partially-ground clauses with one or more facts
 - Edges are proof-transformations
- **Parameters** are learned for each **ground entity and rule**

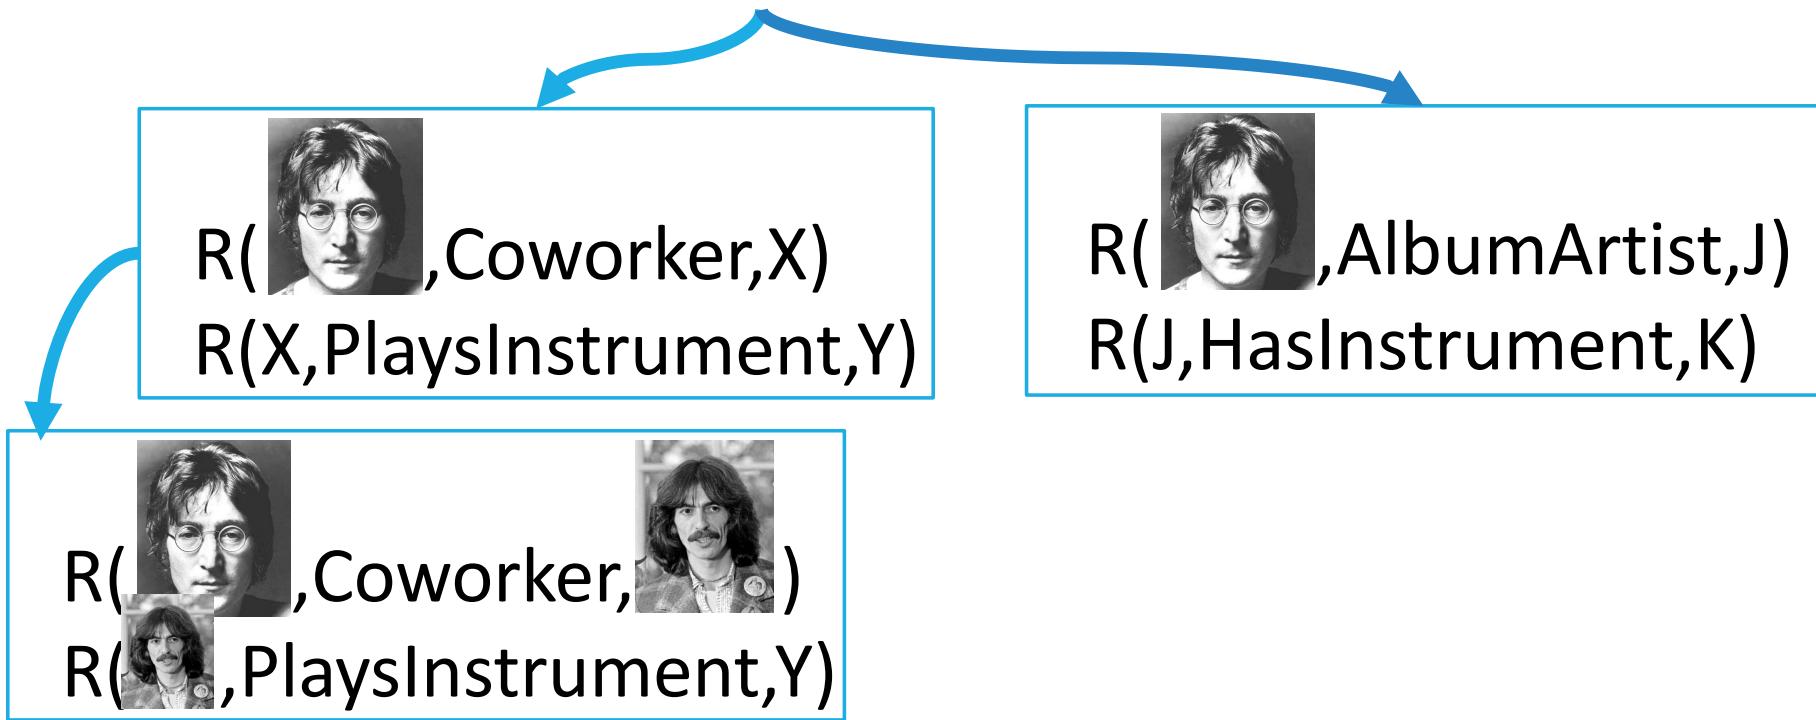
ProPPR-ized PRA example

Query Q: R(Lennon, PlaysInstrument, ?)



ProPPR-ized PRA example

Query Q: R(Lennon, PlaysInstrument, ?)



ProPPR-ized PRA example

Query Q: R(Lennon, PlaysInstrument, ?)

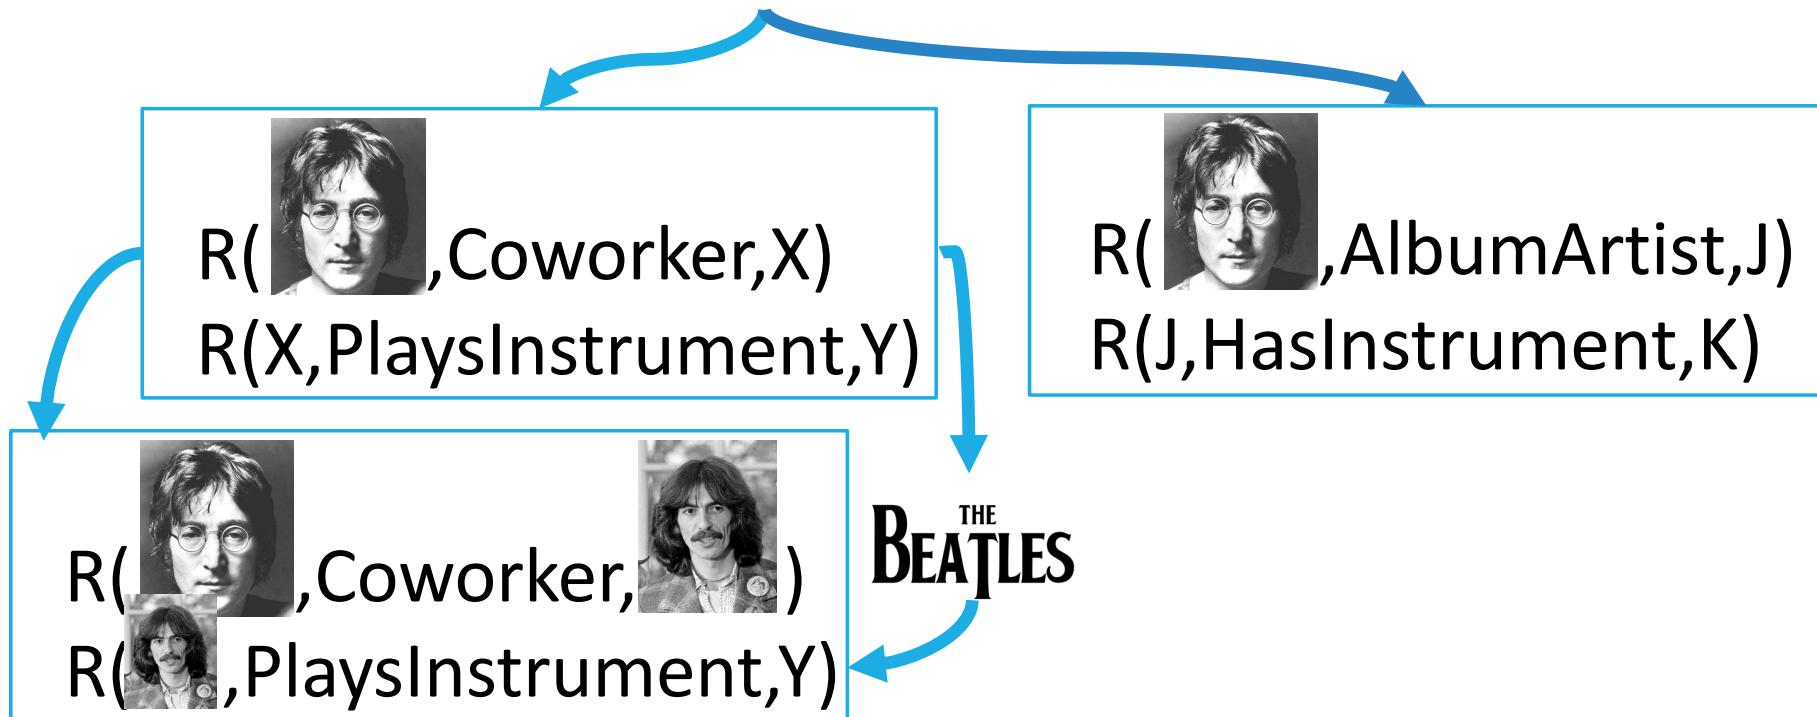
R(,Coworker,X)
R(X,PlaysInstrument,Y)

R(,AlbumArtist,J)
R(J,HasInstrument,K)

R(,Coworker,)
R(,PlaysInstrument,Y)

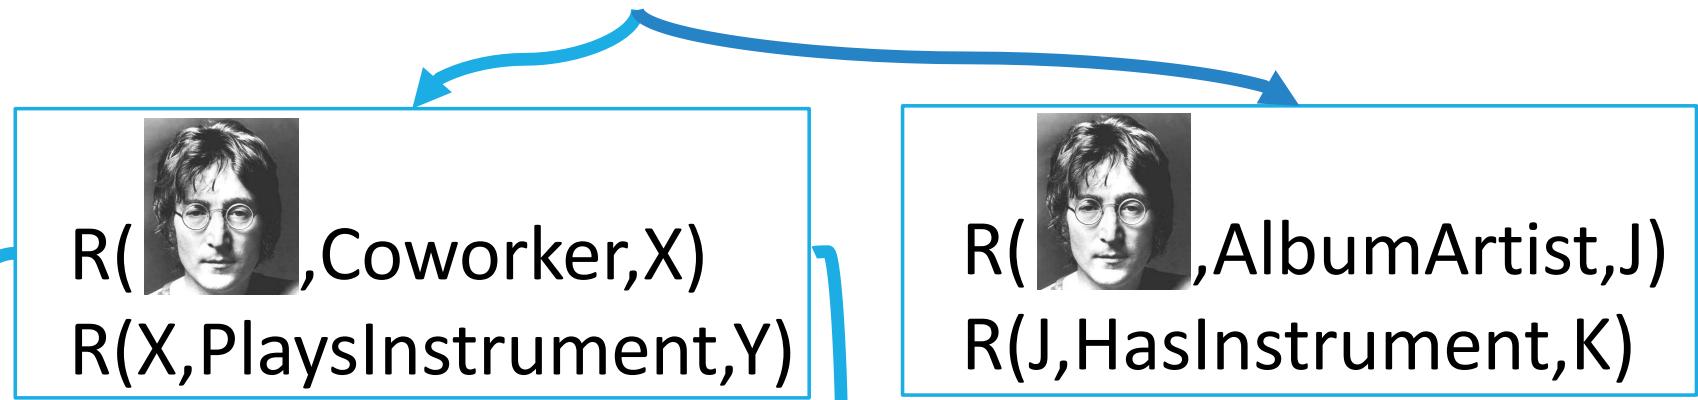
ProPPR-ized PRA example

Query Q: $R(\text{Lennon}, \text{PlaysInstrument}, ?)$

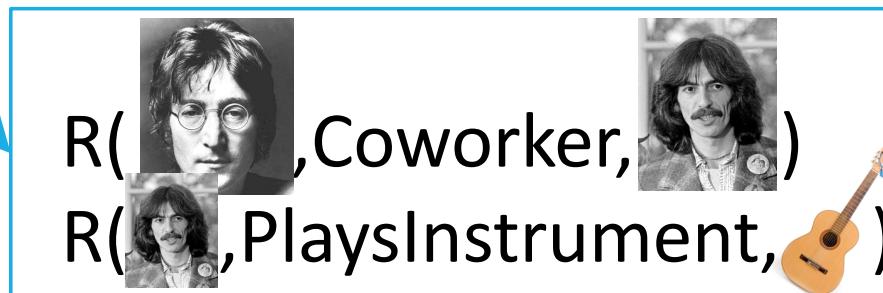


ProPPR-ized PRA example

Query Q: R(Lennon, PlaysInstrument, ?)

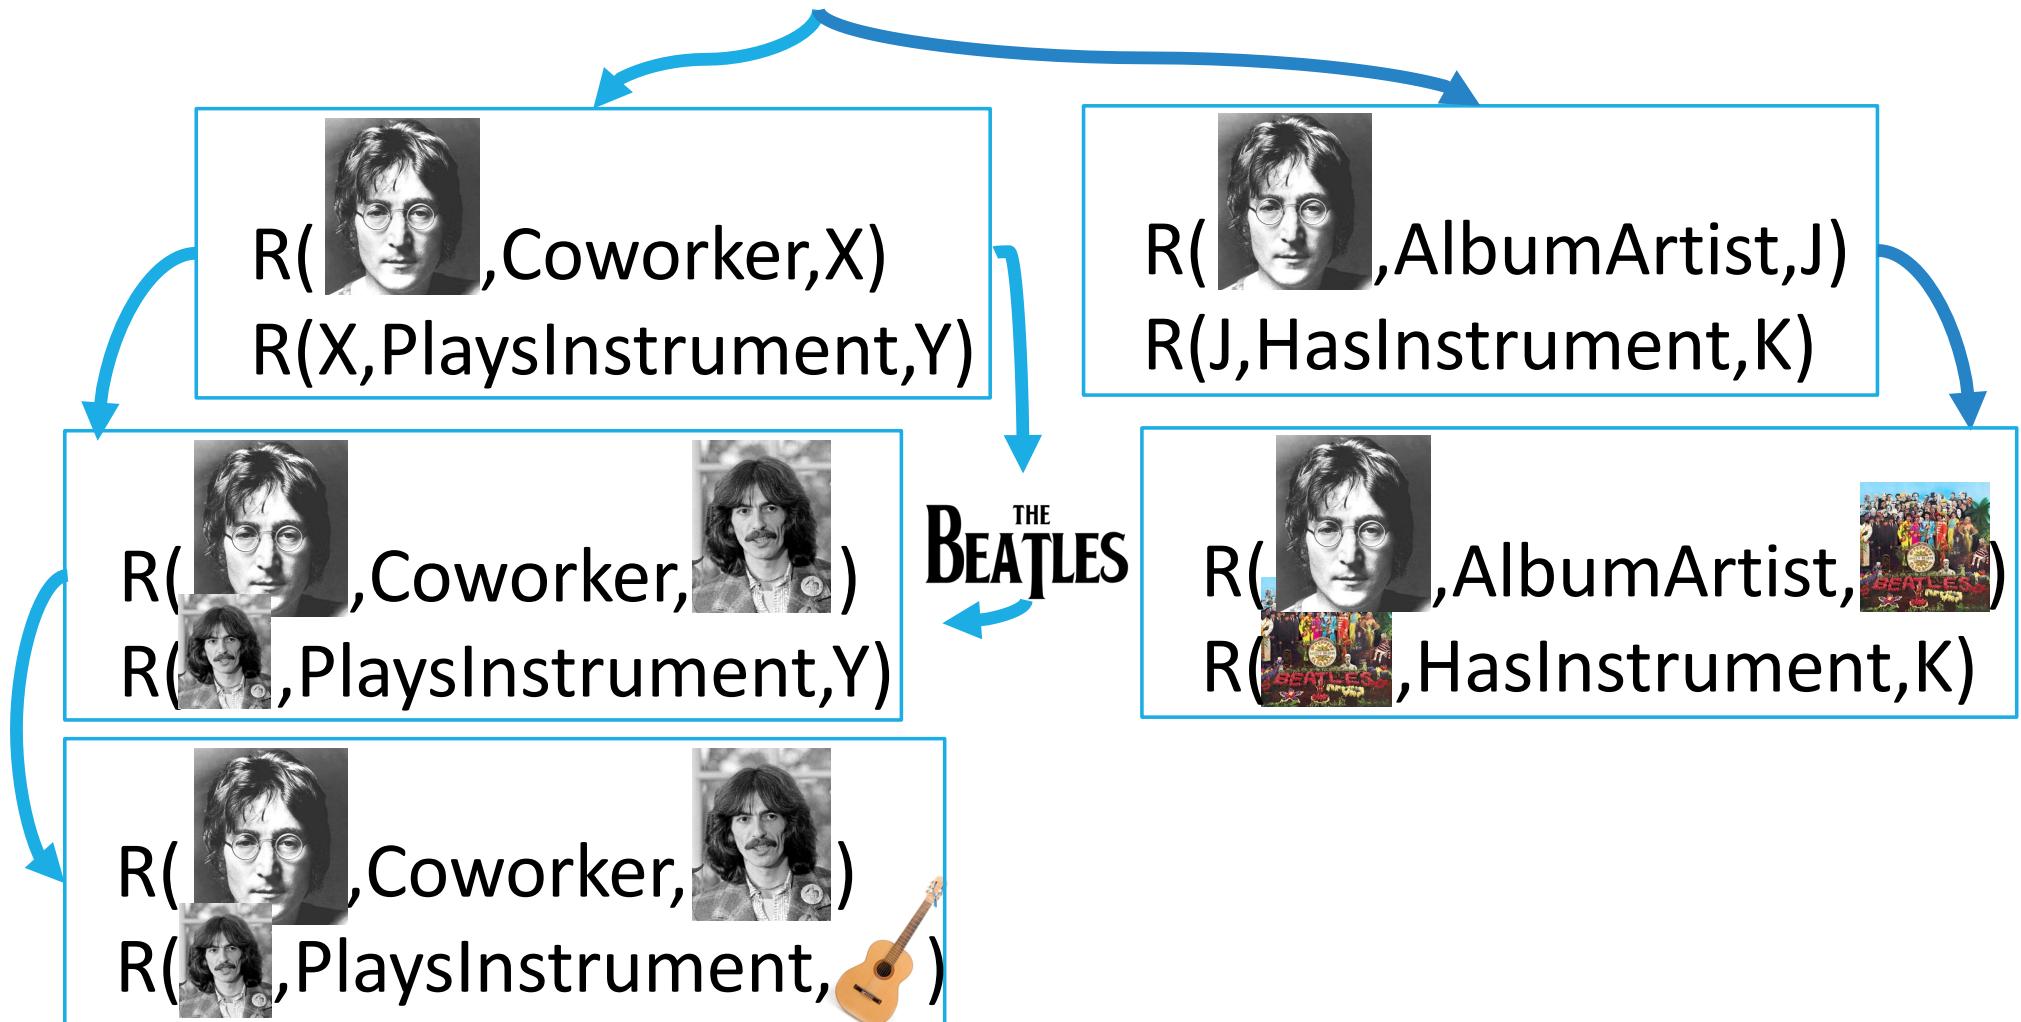


THE
BEATLES



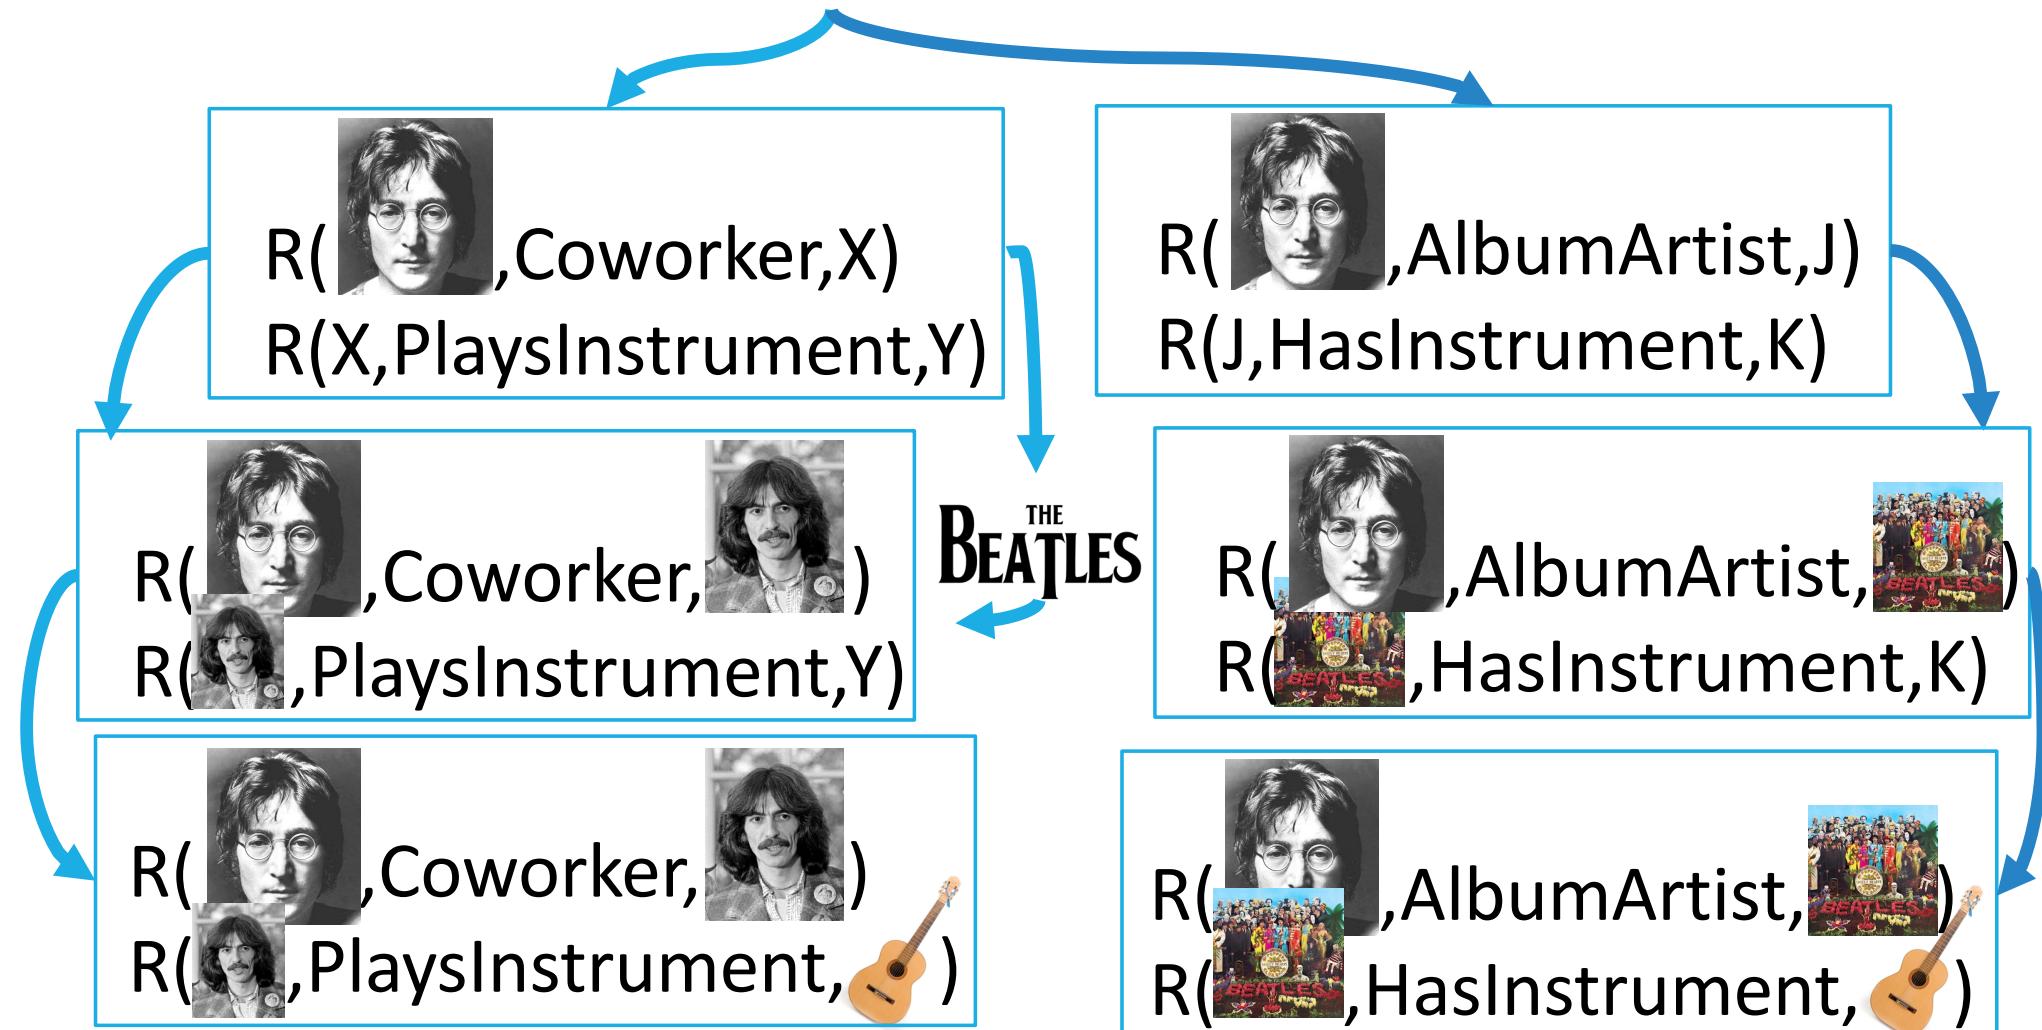
ProPPR-ized PRA example

Query Q: R(Lennon, PlaysInstrument, ?)



ProPPR-ized PRA example

Query Q: R(Lennon, PlaysInstrument, ?)



ProPPR-ized PRA example

Query Q: R(Lennon, PlaysInstrument, ?)



R(John Lennon, Coworker, X)
R(X, PlaysInstrument, ?)



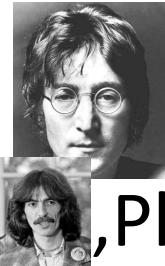
),AlbumArtist,J)
,HasInstrument,K)



R(John Lennon, Coworker, X)
R(X, PlaysInstrument, ?)



),AlbumArtist,
HasInstrument,K)



R(John Lennon, Coworker, X)
R(X, PlaysInstrument, ?)



R(John Lennon, AlbumArtist,
HasInstrument,)



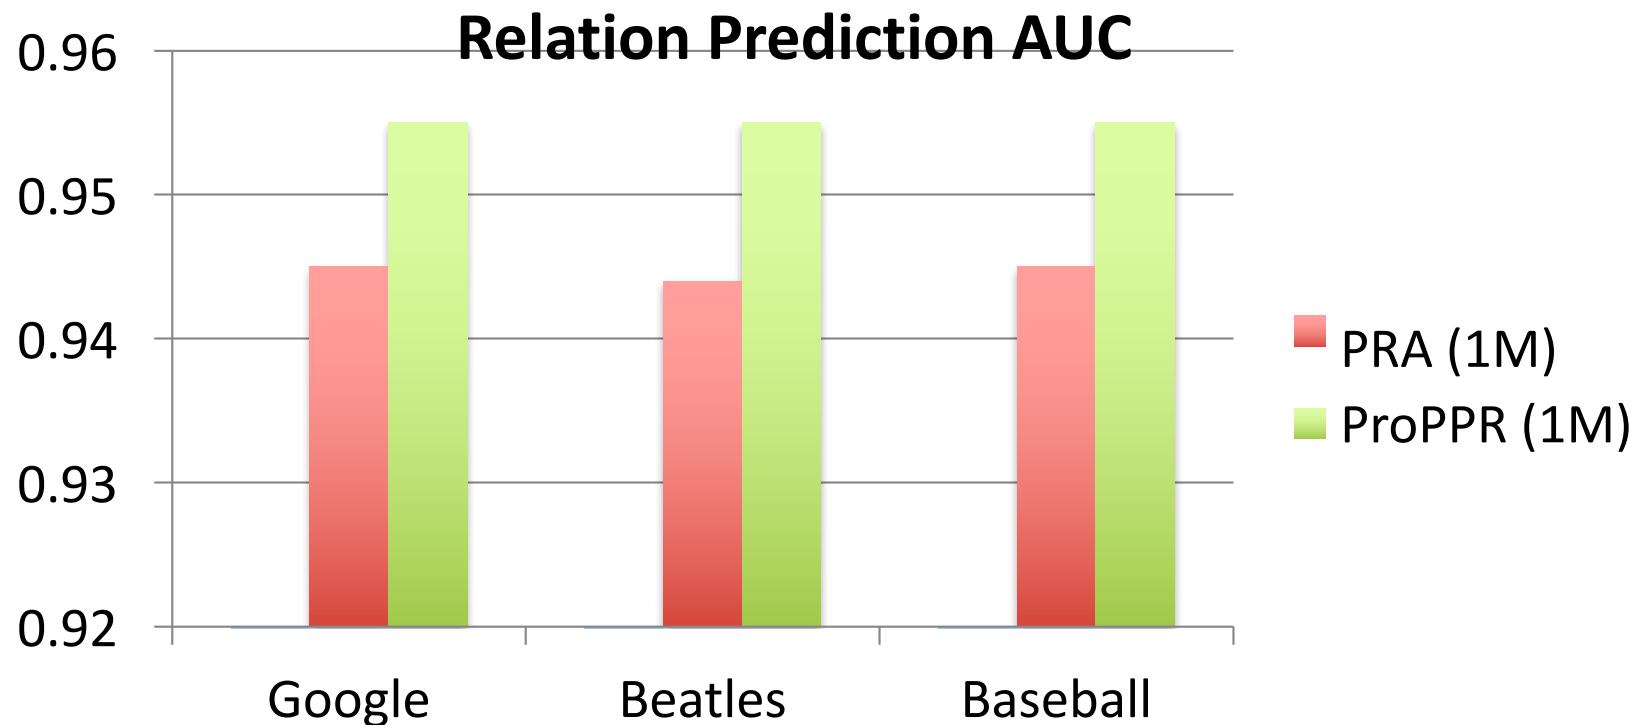
ProPPR in a nutshell

$$\min_{\mathbf{w}} - \left(\sum_{k \in +} \log \mathbf{p}_{\nu_0}[u_+^k] + \sum_{k \in -} \log(1 - \mathbf{p}_{\nu_0}[u_-^k]) \right) + \mu \|\mathbf{w}\|_2^2$$

- Input: queries, positive answers, negative answers
- Goal: $\mathbf{p}_{\nu_0}[u_+^k] \geq \mathbf{p}_{\nu_0}[u_-^k]$ (page rank from RW)
- Learn: random walk weights
- Train via stochastic gradient descent

Results from PRA and ProPPR

- Task:
 - 1M extractions for 3 domains;
 - ~100s of training queries
 - ~1000s of test queries
 - AUC of extractions alone is 0.7



Random Walks: Pros/Cons

BENEFITS

- KG query estimation independent of KG size
- Model training produces interpretable, logical rules
- Robust to noisy extractions through probabilistic form

DRAWBACKS

- Full KG completion task inefficient
- Training data difficult to obtain at scale
- Input must follow probabilistic semantics

Two classes of Probabilistic Models

GRAPHICAL MODELS

- Possible facts in KG are variables
- Logical rules relate facts
- Probability \propto satisfied rules
- Universally-quantified

RANDOM WALK METHODS

- Possible facts posed as queries
- Random walks of the KG constitute “proofs”
- Probability \propto path lengths/transitions
- Locally grounded